

Time series analysis

Article • 05/23/2023

Cloud services and IoT devices generate telemetry data that can be used to gain insights such as monitoring service health, physical production processes, and usage trends.

Performing time series analysis is one way to identify deviations in the pattern of these metrics compared to their typical baseline pattern.

Kusto Query Language (KQL) contains native support for creation, manipulation, and analysis of multiple time series. In this article, learn how KQL is used to create and analyze thousands of time series in seconds, enabling near real-time monitoring solutions and workflows.

Time series creation

In this section, we'll create a large set of regular time series simply and intuitively using the `make-series` operator, and fill-in missing values as needed. The first step in time series analysis is to partition and transform the original telemetry table to a set of time series. The table usually contains a timestamp column, contextual dimensions, and optional metrics. The dimensions are used to partition the data. The goal is to create thousands of time series per partition at regular time intervals.

The input table `demo_make_series1` contains 600K records of arbitrary web service traffic. Use the following command to sample 10 records:

[Run the query](#)

Kusto

```
demo_make_series1 | take 10
```

The resulting table contains a timestamp column, three contextual dimensions columns, and no metrics:

TimeStamp	BrowserVer	OsVer	Country/Region
2016-08-25 09:12:35.4020000	Chrome 51.0	Windows 7	United Kingdom
2016-08-25 09:12:41.1120000	Chrome 52.0	Windows 10	
2016-08-25 09:12:46.2300000	Chrome 52.0	Windows 7	United Kingdom

TimeStamp	BrowserVer	OsVer	Country/Region
2016-08-25 09:12:46.5100000	Chrome 52.0	Windows 10	United Kingdom
2016-08-25 09:12:46.5570000	Chrome 52.0	Windows 10	Republic of Lithuania
2016-08-25 09:12:47.0470000	Chrome 52.0	Windows 8.1	India
2016-08-25 09:12:51.3600000	Chrome 52.0	Windows 10	United Kingdom
2016-08-25 09:12:51.6930000	Chrome 52.0	Windows 7	Netherlands
2016-08-25 09:12:56.4240000	Chrome 52.0	Windows 10	United Kingdom
2016-08-25 09:13:08.7230000	Chrome 52.0	Windows 10	India

Since there are no metrics, we can only build a set of time series representing the traffic count itself, partitioned by OS using the following query:

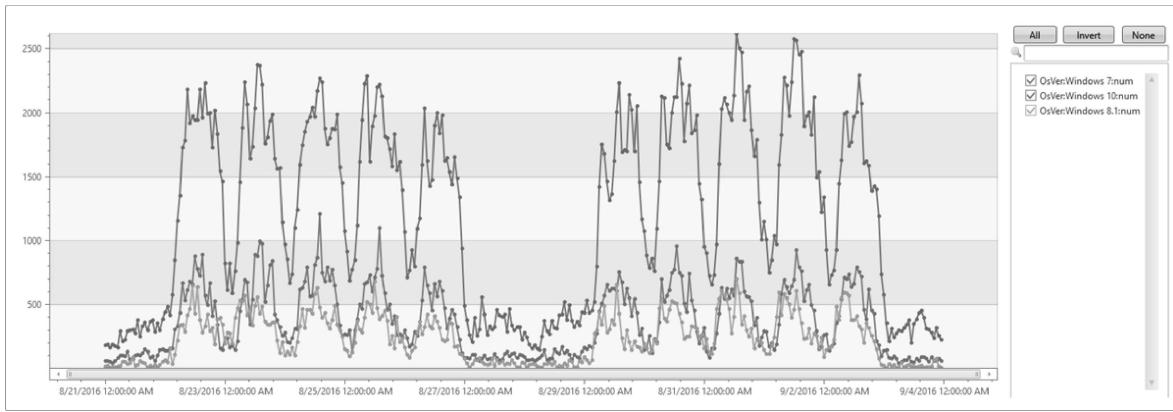
Run the query

Kusto

```
let min_t = toscalar(demo_make_series1 | summarize min(TimeStamp));
let max_t = toscalar(demo_make_series1 | summarize max(TimeStamp));
demo_make_series1
| make-series num=count() default=0 on TimeStamp from min_t to max_t step 1h
by OsVer
| render timechart
```

- Use the make-series operator to create a set of three time series, where:
 - `num=count()`: time series of traffic
 - `from min_t to max_t step 1h`: time series is created in 1-hour bins in the time range (oldest and newest timestamps of table records)
 - `default=0`: specify fill method for missing bins to create regular time series.
Alternatively use `series_fill_const()`, `series_fill_forward()`, `series_fill_backward()` and `series_fill_linear()` for changes
 - `by OsVer`: partition by OS
- The actual time series data structure is a numeric array of the aggregated value per each time bin. We use `render timechart` for visualization.

In the table above, we have three partitions. We can create a separate time series: Windows 10 (red), 7 (blue) and 8.1 (green) for each OS version as seen in the graph:



Time series analysis functions

In this section, we'll perform typical series processing functions. Once a set of time series is created, KQL supports a growing list of functions to process and analyze them. We'll describe a few representative functions for processing and analyzing time series.

Filtering

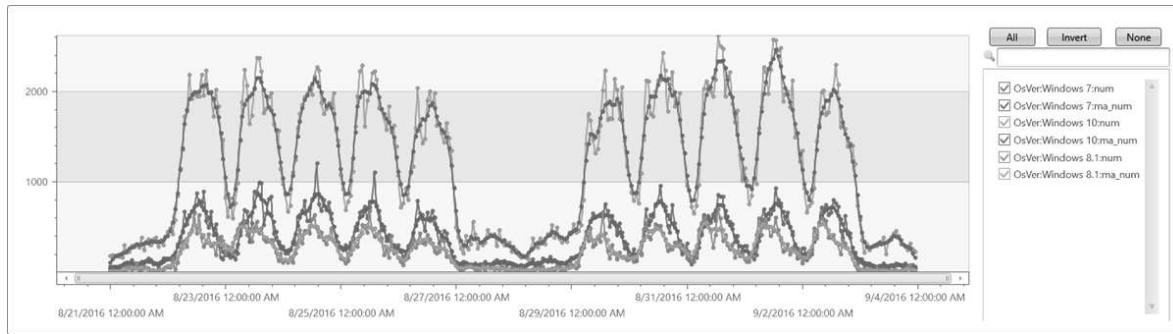
Filtering is a common practice in signal processing and useful for time series processing tasks (for example, smooth a noisy signal, change detection).

- There are two generic filtering functions:
 - `series_fir()`: Applying FIR filter. Used for simple calculation of moving average and differentiation of the time series for change detection.
 - `series_iir()`: Applying IIR filter. Used for exponential smoothing and cumulative sum.
- Extend the time series set by adding a new moving average series of size 5 bins (named *ma_num*) to the query:

Run the query

Kusto

```
let min_t = toscalar(demo_make_series1 | summarize min(TimeStamp));
let max_t = toscalar(demo_make_series1 | summarize max(TimeStamp));
demo_make_series1
| make-series num=count() default=0 on TimeStamp from min_t to max_t step 1h
by OsVer
| extend ma_num=series_fir(num, repeat(1, 5), true, true)
| render timechart
```



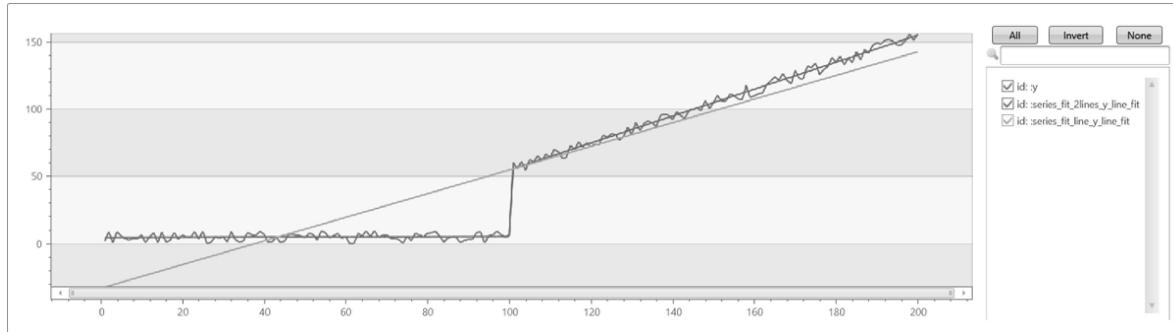
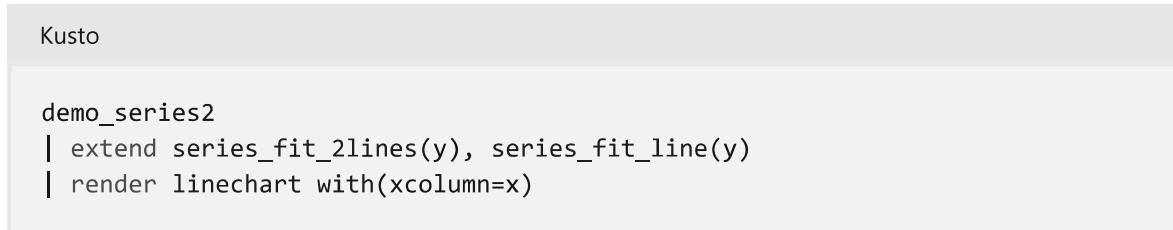
Regression analysis

ADX supports segmented linear regression analysis to estimate the trend of the time series.

- Use `series_fit_line()` to fit the best line to a time series for general trend detection.
- Use `series_fit_2lines()` to detect trend changes, relative to the baseline, that are useful in monitoring scenarios.

Example of `series_fit_line()` and `series_fit_2lines()` functions in a time series query:

Run the query



- Blue: original time series
- Green: fitted line
- Red: two fitted lines

! Note

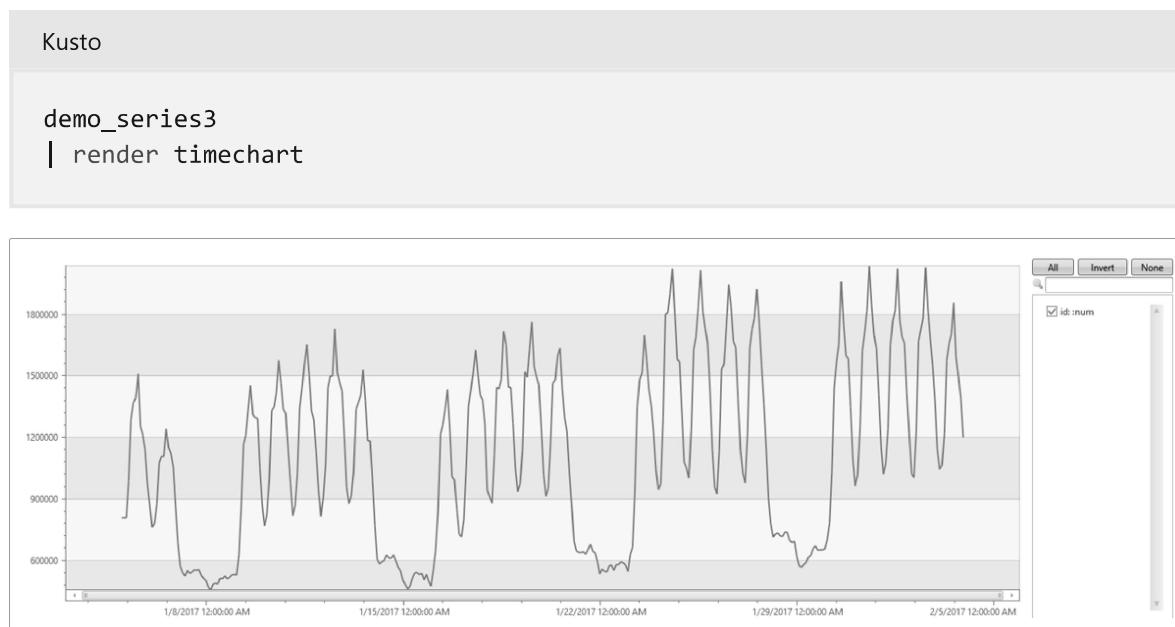
The function accurately detected the jump (level change) point.

Seasonality detection

Many metrics follow seasonal (periodic) patterns. User traffic of cloud services usually contains daily and weekly patterns that are highest around the middle of the business day and lowest at night and over the weekend. IoT sensors measure in periodic intervals. Physical measurements such as temperature, pressure, or humidity may also show seasonal behavior.

The following example applies seasonality detection on one month traffic of a web service (2-hour bins):

Run the query

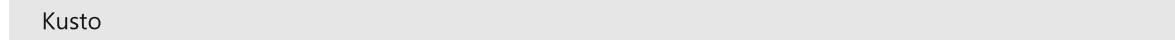


- Use `series_periods_detect()` to automatically detect the periods in the time series.
- Use `series_periods_validate()` if we know that a metric should have specific distinct period(s) and we want to verify that they exist.

! Note

It's an anomaly if specific distinct periods don't exist

Run the query



```

demo_series3
| project (periods, scores) = series_periods_detect(num, 0., 14d/2h, 2) //to
detect the periods in the time series
| mv-expand periods, scores
| extend days=2h*todouble(periods)/1d

```

periods	scores	days
84	0.820622786055595	7
12	0.764601405803502	1

The function detects daily and weekly seasonality. The daily scores less than the weekly because weekend days are different from weekdays.

Element-wise functions

Arithmetic and logical operations can be done on a time series. Using `series_subtract()` we can calculate a residual time series, that is, the difference between original raw metric and a smoothed one, and look for anomalies in the residual signal:

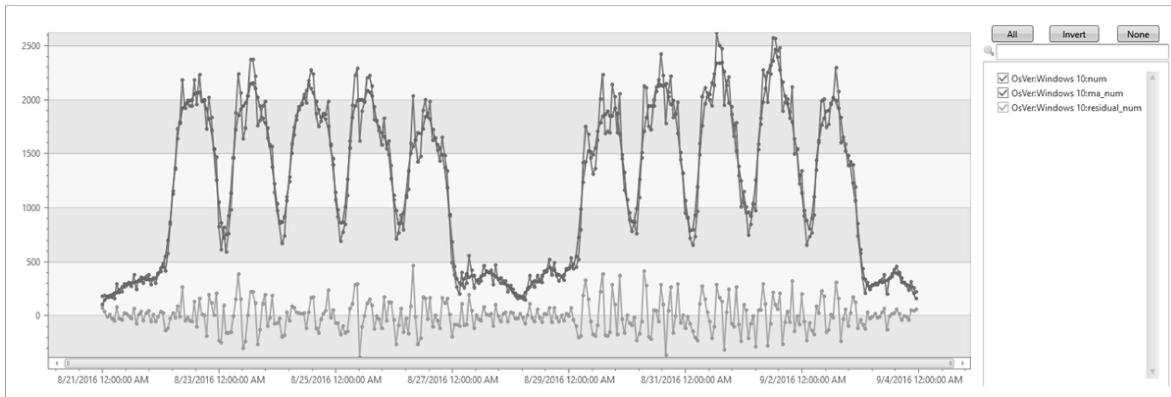
[Run the query](#)

Kusto

```

let min_t = toscalar(demo_make_series1 | summarize min(TimeStamp));
let max_t = toscalar(demo_make_series1 | summarize max(TimeStamp));
demo_make_series1
| make-series num=count() default=0 on TimeStamp in from min_t to max_t step
1h by OsVer
| extend ma_num=series_fir(num, repeat(1, 5), true, true)
| extend residual_num=series_subtract(num, ma_num) //to calculate residual
time series
| where OsVer == "Windows 10"    // filter on Win 10 to visualize a cleaner
chart
| render timechart

```



- Blue: original time series
- Red: smoothed time series
- Green: residual time series

Time series workflow at scale

The example below shows how these functions can run at scale on thousands of time series in seconds for anomaly detection. To see a few sample telemetry records of a DB service's read count metric over four days run the following query:

[Run the query](#)

Kusto

```
demo_many_series1
| take 4
```

TIMESTAMP	Loc	anonOp	DB	DataRead
2016-09-11 21:00:00.0000000	Loc 9	5117853934049630089	262	0
2016-09-11 21:00:00.0000000	Loc 9	5117853934049630089	241	0
2016-09-11 21:00:00.0000000	Loc 9	-865998331941149874	262	279862
2016-09-11 21:00:00.0000000	Loc 9	371921734563783410	255	0

And simple statistics:

[Run the query](#)

Kusto

```
demo_many_series1
| summarize num=count(), min_t=min(TIMESTAMP), max_t=max(TIMESTAMP)
```

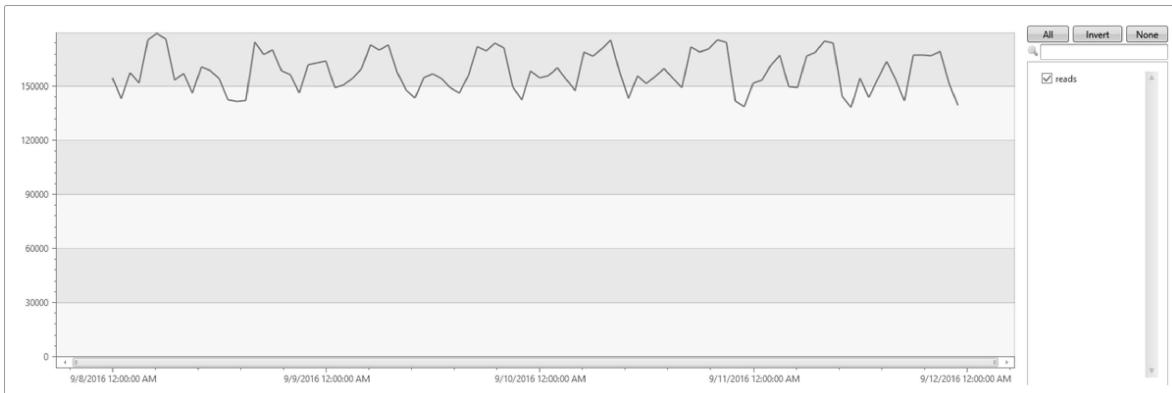
num	min_t	max_t
2177472	2016-09-08 00:00:00.0000000	2016-09-11 23:00:00.0000000

Building a time series in 1-hour bins of the read metric (total four days * 24 hours = 96 points), results in normal pattern fluctuation:

Run the query

Kusto

```
let min_t = toscalar(demo_many_series1 | summarize min(TIMESTAMP));
let max_t = toscalar(demo_many_series1 | summarize max(TIMESTAMP));
demo_many_series1
| make-series reads=avg(DataRead) on TIMESTAMP from min_t to max_t step 1h
| render timechart with(ymin=0)
```



The above behavior is misleading, since the single normal time series is aggregated from thousands of different instances that may have abnormal patterns. Therefore, we create a time series per instance. An instance is defined by Loc (location), anonOp (operation), and DB (specific machine).

How many time series can we create?

Run the query

Kusto

```
demo_many_series1
| summarize by Loc, Op, DB
| count
```

Count

Count

18339

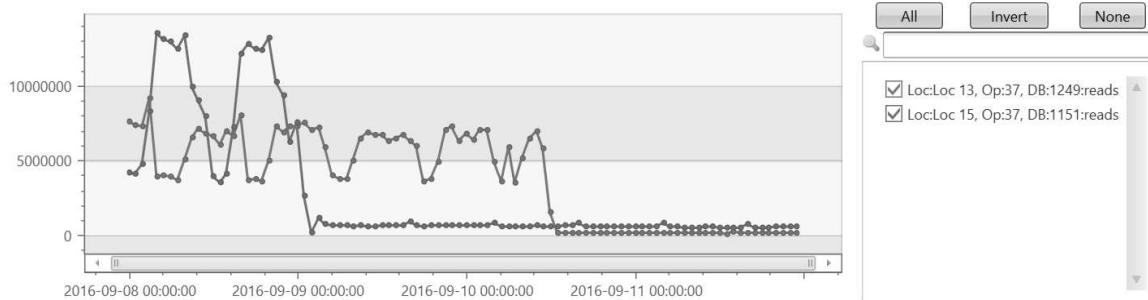
Now, we're going to create a set of 18339 time series of the read count metric. We add the `by` clause to the make-series statement, apply linear regression, and select the top two time series that had the most significant decreasing trend:

Run the query

Kusto

```
let min_t = toscalar(demo_many_series1 | summarize min(TIMESTAMP));
let max_t = toscalar(demo_many_series1 | summarize max(TIMESTAMP));
demo_many_series1
| make-series reads=avg(DataRead) on TIMESTAMP from min_t to max_t step 1h
by Loc, Op, DB
| extend (rsquare, slope) = series_fit_line(reads)
| top 2 by slope asc
| render timechart with(title='Service Traffic Outage for 2 instances (out
of 18339)')
```

Service Traffic Outage for 2 instances (out of 18339)



Display the instances:

Run the query

Kusto

```
let min_t = toscalar(demo_many_series1 | summarize min(TIMESTAMP));
let max_t = toscalar(demo_many_series1 | summarize max(TIMESTAMP));
demo_many_series1
| make-series reads=avg(DataRead) on TIMESTAMP from min_t to max_t step 1h
by Loc, Op, DB
| extend (rsquare, slope) = series_fit_line(reads)
| top 2 by slope asc
| project Loc, Op, DB, slope
```

Loc	Op	DB	slope
Loc 15	37	1151	-102743.910227889
Loc 13	37	1249	-86303.2334644601

In less than two minutes, close to 20,000 time series were analyzed and two abnormal time series in which the read count suddenly dropped were detected.

These advanced capabilities combined with fast performance supply a unique and powerful solution for time series analysis.

Next steps

- Learn about Anomaly detection and forecasting with KQL.
- Learn about Machine learning capabilities with KQL.

Feedback

Was this page helpful?



Provide product feedback | Get help at Microsoft Q&A

Anomaly detection and forecasting

Article • 05/23/2023

Cloud services and IoT devices generate telemetry data that can be used to gain insights such as monitoring service health, physical production processes, and usage trends.

Performing time series analysis is one way to identify deviations in the pattern of these metrics compared to their typical baseline pattern.

Kusto Query Language (KQL) contains native support for creation, manipulation, and analysis of multiple time series. With KQL, you can create and analyze thousands of time series in seconds, enabling near real time monitoring solutions and workflows.

This article details time series anomaly detection and forecasting capabilities of KQL. The applicable time series functions are based on a robust well-known decomposition model, where each original time series is decomposed into seasonal, trend, and residual components. Anomalies are detected by outliers on the residual component, while forecasting is done by extrapolating the seasonal and trend components. The KQL implementation significantly enhances the basic decomposition model by automatic seasonality detection, robust outlier analysis, and vectorized implementation to process thousands of time series in seconds.

Prerequisites

- A Microsoft account or an Azure Active Directory user identity. An Azure subscription isn't required.
- Read Time series analysis for an overview of time series capabilities.

Time series decomposition model

The KQL native implementation for time series prediction and anomaly detection uses a well-known decomposition model. This model is applied to time series of metrics expected to manifest periodic and trend behavior, such as service traffic, component heartbeats, and IoT periodic measurements to forecast future metric values and detect anomalous ones. The assumption of this regression process is that other than the previously known seasonal and trend behavior, the time series is randomly distributed. You can then forecast future metric values from the seasonal and trend components, collectively named baseline, and ignore the residual part. You can also detect anomalous values based on outlier analysis using only the residual portion. To create a decomposition model, use the function `series_decompose()`. The `series_decompose()`

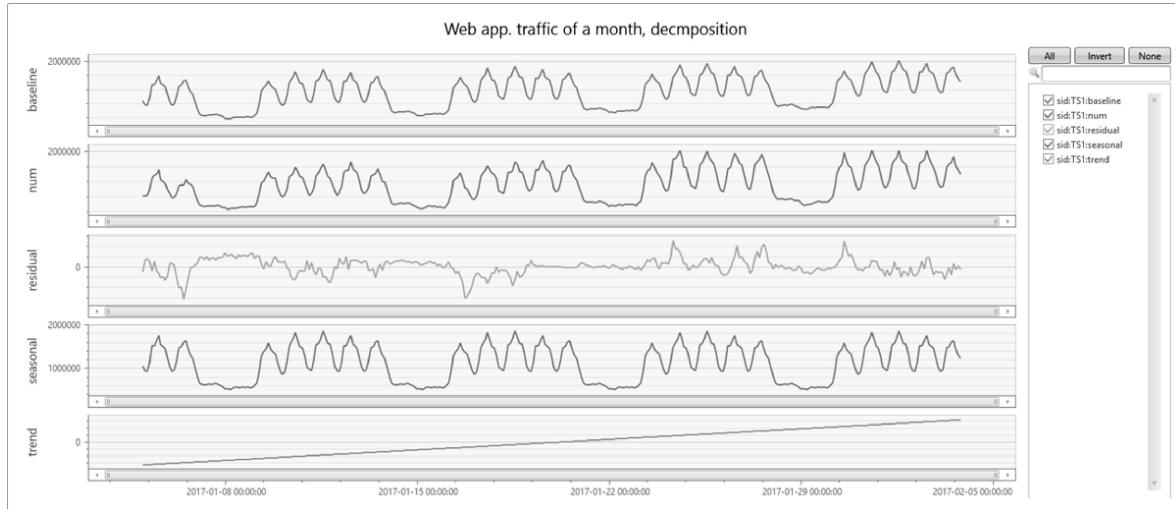
function takes a set of time series and automatically decomposes each time series to its seasonal, trend, residual, and baseline components.

For example, you can decompose traffic of an internal web service by using the following query:

Run the query

Kusto

```
let min_t = datetime(2017-01-05);
let max_t = datetime(2017-02-03 22:00);
let dt = 2h;
demo_make_series2
| make-series num=avg(num) on TimeStamp from min_t to max_t step dt by sid
| where sid == 'TS1'    // select a single time series for a cleaner
visualisation
| extend (baseline, seasonal, trend, residual) = series_decompose(num, -1,
'linefit') // decomposition of a set of time series to seasonal, trend,
residual, and baseline (seasonal+trend)
| render timechart with(title='Web app. traffic of a month, decomposition',
ysplit=panels)
```



- The original time series is labeled **num** (in red).
- The process starts by auto detection of the seasonality by using the function `series_periods_detect()` and extracts the **seasonal** pattern (in purple).
- The seasonal pattern is subtracted from the original time series and a linear regression is run using the function `series_fit_line()` to find the **trend** component (in light blue).
- The function subtracts the trend and the remainder is the **residual** component (in green).
- Finally, the function adds the seasonal and trend components to generate the **baseline** (in blue).

Time series anomaly detection

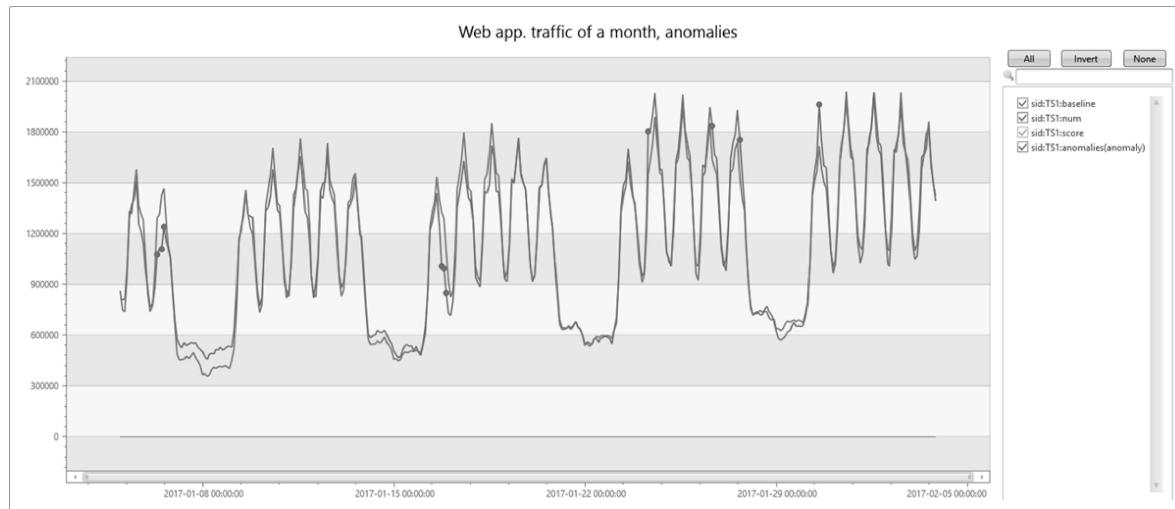
The function `series_decompose_anomalies()` finds anomalous points on a set of time series. This function calls `series_decompose()` to build the decomposition model and then runs `series_outliers()` on the residual component. `series_outliers()` calculates anomaly scores for each point of the residual component using Tukey's fence test. Anomaly scores above 1.5 or below -1.5 indicate a mild anomaly rise or decline respectively. Anomaly scores above 3.0 or below -3.0 indicate a strong anomaly.

The following query allows you to detect anomalies in internal web service traffic:

Run the query

Kusto

```
let min_t = datetime(2017-01-05);
let max_t = datetime(2017-02-03 22:00);
let dt = 2h;
demo_make_series2
| make-series num=avg(num) on TimeStamp from min_t to max_t step dt by sid
| where sid == 'TS1'    //  select a single time series for a cleaner
visualization
| extend (anomalies, score, baseline) = series_decompose_anomalies(num, 1.5,
-1, 'linefit')
| render anomalychart with(anomalycolumns=anomalies, title='Web app. traffic
of a month, anomalies') //use "| render anomalychart with
anomalycolumns=anomalies" to render the anomalies as bold points on the
series charts.
```



- The original time series (in red).
- The baseline (seasonal + trend) component (in blue).
- The anomalous points (in purple) on top of the original time series. The anomalous points significantly deviate from the expected baseline values.

Time series forecasting

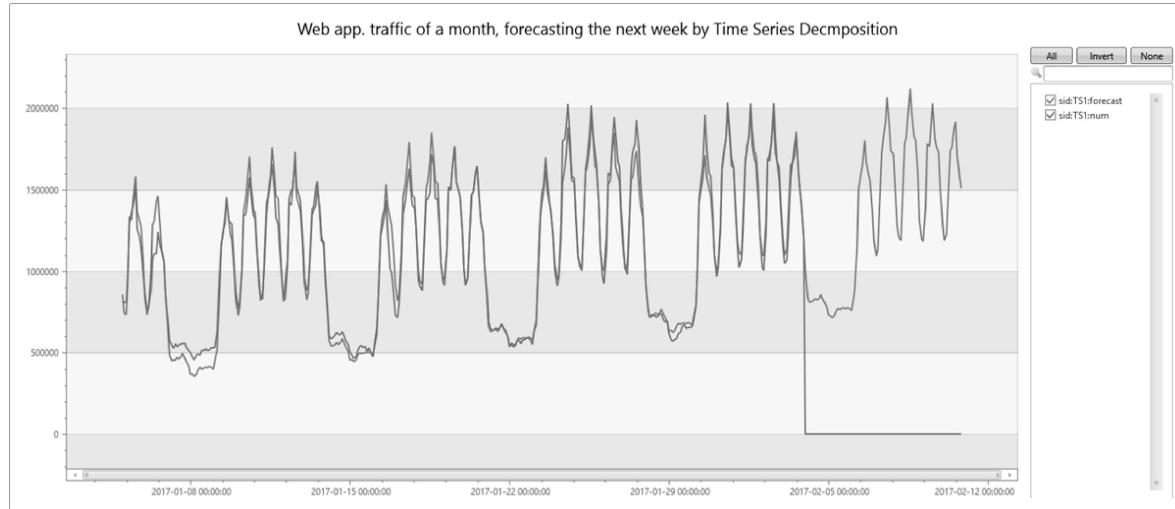
The function `series_decompose_forecast()` predicts future values of a set of time series. This function calls `series_decompose()` to build the decomposition model and then, for each time series, extrapolates the baseline component into the future.

The following query allows you to predict next week's web service traffic:

[Run the query](#)

Kusto

```
let min_t = datetime(2017-01-05);
let max_t = datetime(2017-02-03 22:00);
let dt = 2h;
let horizon=7d;
demo_make_series2
| make-series num=avg(num) on TimeStamp from min_t to max_t+horizon step dt
by sid
| where sid == 'TS1'    //  select a single time series for a cleaner
visualization
| extend forecast = series_decompose_forecast(num, toint(horizon/dt))
| render timechart with(title='Web app. traffic of a month, forecasting the
next week by Time Series Decomposition')
```



- Original metric (in red). Future values are missing and set to 0, by default.
- Extrapolate the baseline component (in blue) to predict next week's values.

Scalability

Kusto Query Language syntax enables a single call to process multiple time series. Its unique optimized implementation allows for fast performance, which is critical for

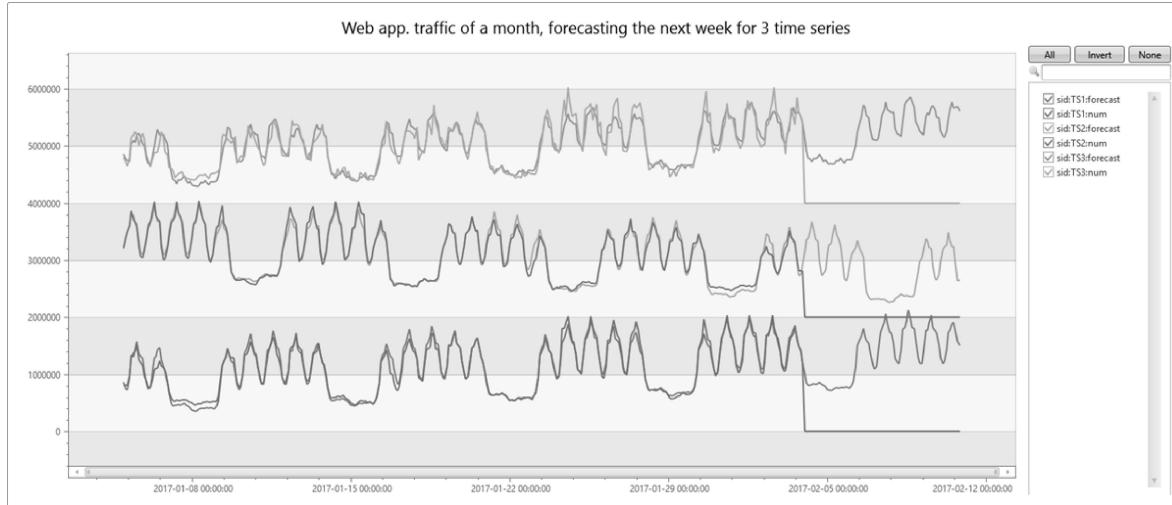
effective anomaly detection and forecasting when monitoring thousands of counters in near real-time scenarios.

The following query shows the processing of three time series simultaneously:

[Run the query](#)

```
Kusto

let min_t = datetime(2017-01-05);
let max_t = datetime(2017-02-03 22:00);
let dt = 2h;
let horizon=7d;
demo_make_series2
| make-series num=avg(num) on TimeStamp from min_t to max_t+horizon step dt
by sid
| extend offset=case(sid=='TS3', 4000000, sid=='TS2', 2000000, 0) // add
artificial offset for easy visualization of multiple time series
| extend num=series_add(num, offset)
| extend forecast = series_decompose_forecast(num, toint(horizon/dt))
| render timechart with(title='Web app. traffic of a month, forecasting the
next week for 3 time series')
```



Summary

This document details native KQL functions for time series anomaly detection and forecasting. Each original time series is decomposed into seasonal, trend and residual components for detecting anomalies and/or forecasting. These functionalities can be used for near real-time monitoring scenarios, such as fault detection, predictive maintenance, and demand and load forecasting.

Next steps

Learn about Machine learning capabilities with KQL.

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

Machine learning capability

Article • 05/23/2023

Kusto Query Language (KQL) has built-in anomaly detection and forecasting functions to check for anomalous behavior. Once such a pattern is detected, a Root Cause Analysis (RCA) can be run to mitigate or resolve the anomaly.

The diagnosis process is complex and lengthy, and done by domain experts. The process includes:

- Fetching and joining additional data from different sources for the same time frame
- Looking for changes in the distribution of values on multiple dimensions
- Charting additional variables
- Other techniques based on domain knowledge and intuition

Since these diagnosis scenarios are common, machine learning plugins are available to make the diagnosis phase easier, and shorten the duration of the RCA.

All three of the following Machine Learning plugins implement clustering algorithms: autocluster, basket, and diffpatterns. The `autocluster` and `basket` plugins cluster a single record set, and the `diffpatterns` plugin clusters the differences between two record sets.

Clustering a single record set

A common scenario includes a data set selected by a specific criteria such as:

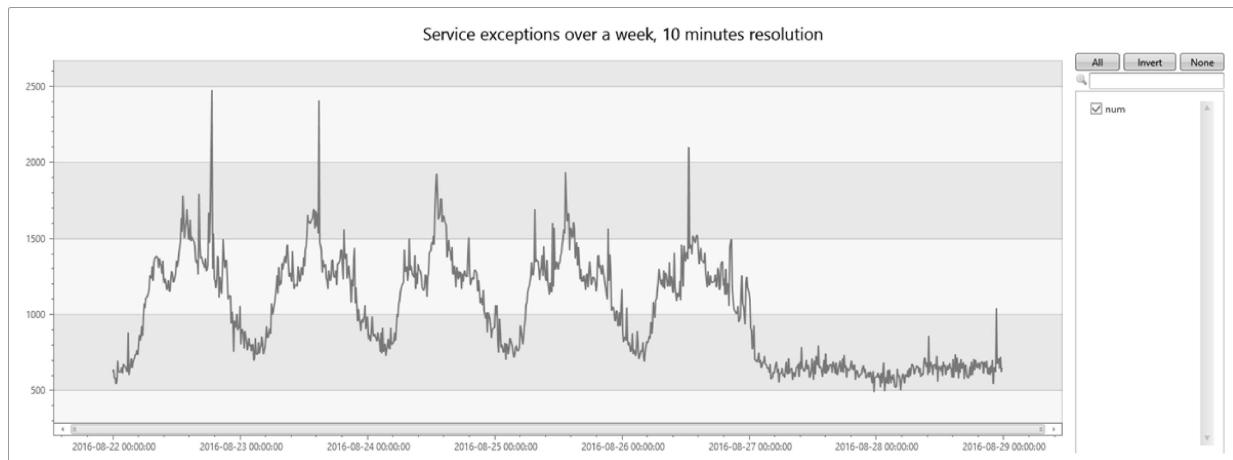
- Time window that shows anomalous behavior
- High temperature device readings
- Long duration commands
- Top spending users You want a fast and easy way to find common patterns (segments) in the data. Patterns are a subset of the data set whose records share the same values over multiple dimensions (categorical columns).

The following query builds and shows a time series of service exceptions over the period of a week, in ten-minute bins:

Run the query

Kusto

```
let min_t = toscalar(demo_clustering1 | summarize min(PreciseTimeStamp));
let max_t = toscalar(demo_clustering1 | summarize max(PreciseTimeStamp));
demo_clustering1
| make-series num=count() on PreciseTimeStamp from min_t to max_t step 10m
| render timechart with(title="Service exceptions over a week, 10 minutes resolution")
```



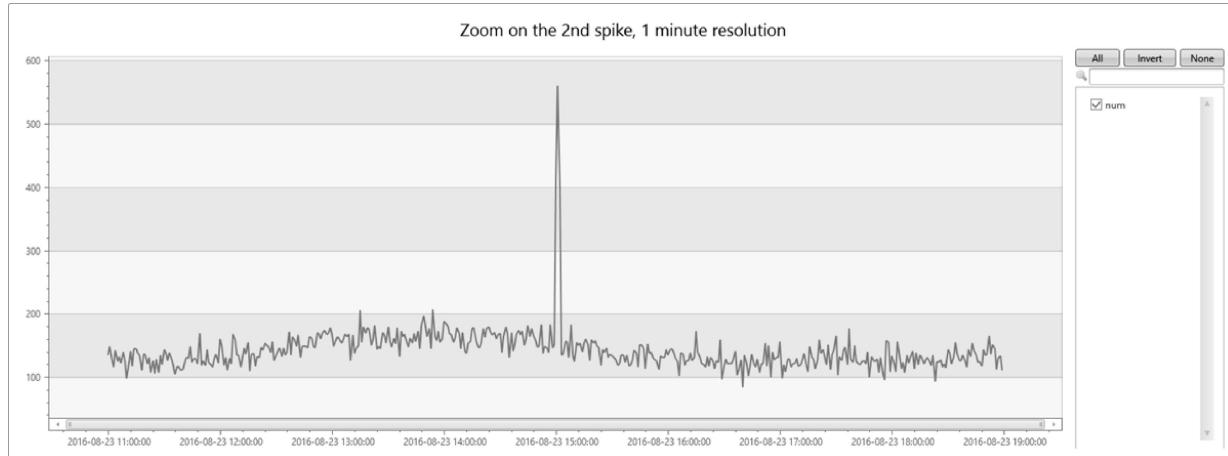
The service exception count correlates with the overall service traffic. You can clearly see the daily pattern for business days, Monday to Friday. There's a rise in service exception counts at mid-day, and drops in counts during the night. Flat low counts are visible over the weekend. Exception spikes can be detected using time series anomaly detection.

The second spike in the data occurs on Tuesday afternoon. The following query is used to further diagnose and verify whether it's a sharp spike. The query redraws the chart around the spike in a higher resolution of eight hours in one-minute bins. You can then study its borders.

Run the query

Kusto

```
let min_t=datetime(2016-08-23 11:00);
demo_clustering1
| make-series num=count() on PreciseTimeStamp from min_t to min_t+8h step 1m
| render timechart with(title="Zoom on the 2nd spike, 1 minute resolution")
```



You'll see a narrow two-minute spike from 15:00 to 15:02. In the following query, count the exceptions in this two-minute window:

Run the query

Kusto

```
let min_peak_t=datetime(2016-08-23 15:00);
let max_peak_t=datetime(2016-08-23 15:02);
demo_clustering1
| where PreciseTimeStamp between(min_peak_t..max_peak_t)
| count
```

Count

972

In the following query, sample 20 exceptions out of 972:

Run the query

Kusto

```
let min_peak_t=datetime(2016-08-23 15:00);
let max_peak_t=datetime(2016-08-23 15:02);
demo_clustering1
| where PreciseTimeStamp between(min_peak_t..max_peak_t)
| take 20
```

PreciseTimeStamp	Region	ScaleUnit	DeploymentId	Tracepoint	ServiceHost
2016-08-23 15:00:08.7302460	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	100005	00000000-0000-0000-0000-000000000000
2016-08-23 15:00:09.9496584	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	10007006	8d257da1-7a1c-44f5-9acd-f9e02ff507fd

PreciseTimeStamp	Region	ScaleUnit	DeploymentId	Tracepoint	ServiceHost
2016-08-23 15:00:10.5911748	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	100005	00000000-0000-0000-0000-000000000000
2016-08-23 15:00:12.2957912	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	10007007	f855fcf-ebfe-405d-aaf8-9c5e2e43d862
2016-08-23 15:00:18.5955357	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	10007006	9d390e07-417d-42eb-bebd-793965189a28
2016-08-23 15:00:20.7444854	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	10007006	6e54c1c8-42d3-4e4e-8b79-9bb076ca71f1
2016-08-23 15:00:23.8694999	eus2	su2	89e2f62a73bb4efd8f545aeae40d7e51	36109	19422243-19b9-4d85-9ca6-bc961861d287
2016-08-23 15:00:26.4271786	ncus	su1	e24ef436e02b4823ac5d5b1465a9401e	36109	3271bae4-1c5b-4f73-98ef-cc117e9be914
2016-08-23 15:00:27.8958124	scus	su3	90d3d2fc7ecc430c9621ece335651a01	904498	8cf38575-fca9-48ca-bd7c-21196f6d6765
2016-08-23 15:00:32.9884969	scus	su3	90d3d2fc7ecc430c9621ece335651a01	10007007	d5c7c825-9d46-4ab7-a0c1-8e2ac1d83ddb
2016-08-23 15:00:34.5061623	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	1002110	55a71811-5ec4-497a-a058-140fb0d611ad
2016-08-23 15:00:37.4490273	scus	su3	90d3d2fc7ecc430c9621ece335651a01	10007006	f2ee8254-173c-477d-a1de-4902150ea50d
2016-08-23 15:00:41.2431223	scus	su3	90d3d2fc7ecc430c9621ece335651a01	103200	8cf38575-fca9-48ca-bd7c-21196f6d6765
2016-08-23 15:00:47.2983975	ncus	su1	e24ef436e02b4823ac5d5b1465a9401e	423690590	00000000-0000-0000-0000-000000000000
2016-08-23 15:00:50.5932834	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	10007006	2a41b552-aa19-4987-8cdd-410a3af016ac
2016-08-23 15:00:50.8259021	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	1002110	0d56b8e3-470d-4213-91da-97405f8d005e
2016-08-23 15:00:53.2490731	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	36109	55a71811-5ec4-497a-a058-140fb0d611ad
2016-08-23 15:00:57.0000946	eus2	su2	89e2f62a73bb4efd8f545aeae40d7e51	64038	cb55739e-4afe-46a3-970f-1b49d8ee7564
2016-08-23 15:00:58.2222707	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	10007007	8215dcf6-2de0-42bd-9c90-181c70486c9c
2016-08-23 15:00:59.9382620	scus	su3	90d3d2fc7ecc430c9621ece335651a01	10007006	451e3c4c-0808-4566-a64d-84d85cf30978

Use autocluster() for single record set clustering

Even though there are less than a thousand exceptions, it's still hard to find common segments, since there are multiple values in each column. You can use the autocluster() plugin to instantly extract a short list of common segments and find the interesting clusters within the spike's two minutes, as seen in the following query:

[Run the query](#)

Kusto

```
let min_peak_t=datetime(2016-08-23 15:00);
let max_peak_t=datetime(2016-08-23 15:02);
demo_clustering1
| where PreciseTimeStamp between(min_peak_t..max_peak_t)
| evaluate autocluster()
```

SegmentId	Count	Percent	Region	ScaleUnit	DeploymentId	ServiceHost
0	639	65.7407407407407	eau	su7	b5d1d4df547d4a04ac15885617edba57	e7f60c5d-4944-42b3-922a-92e98a8e7dec
1	94	9.67078189300411	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6	
2	82	8.43621399176955	ncus	su1	e24ef436e02b4823ac5d5b1465a9401e	
3	68	6.99588477366255	scus	su3	90d3d2fc7ecc430c9621ece335651a01	
4	55	5.65843621399177	weu	su4	be1d6d7ac9574cbc9a22cb8ee20f16fc	

You can see from the results above, the most dominant segment contains 65.74% of the total exception records and shares four dimensions. The next segment is much less common. It contains only 9.67% of the records, and shares three dimensions. The other segments are even less common.

Autocluster uses a proprietary algorithm for mining multiple dimensions and extracting interesting segments. "Interesting" means that each segment has significant coverage of both the records set and the features set. The segments are also diverged, meaning that each one is different from the others. One or more of these segments may be relevant for the RCA process. To minimize segment review and assessment, autocluster extracts only a small segment list.

Use basket() for single record set clustering

You can also use the basket() plugin as seen in the following query:

Run the query

Kusto

```
let min_peak_t=datetime(2016-08-23 15:00);
let max_peak_t=datetime(2016-08-23 15:02);
demo_clustering1
| where PreciseTimeStamp between(min_peak_t..max_peak_t)
| evaluate basket()
```

SegmentId	Count	Percent	Region	ScaleUnit	DeploymentId	Tracepoint	ServiceHost
0	639	65.7407407407407	eau	su7	b5d1d4df547d4a04ac15885617edba57		e7f60c5d-4944-42b3-922a-92e98a8e7dec
1	642	66.0493827160494	eau	su7	b5d1d4df547d4a04ac15885617edba57		
2	324	33.3333333333333	eau	su7	b5d1d4df547d4a04ac15885617edba57	0	e7f60c5d-4944-42b3-922a-92e98a8e7dec
3	315	32.4074074074074	eau	su7	b5d1d4df547d4a04ac15885617edba57	16108	e7f60c5d-4944-42b3-922a-92e98a8e7dec
4	328	33.7448559670782				0	
5	94	9.67078189300411	scus	su5	9dbd1b161d5b4779a73cf19a7836ebd6		
6	82	8.43621399176955	ncus	su1	e24ef436e02b4823ac5d5b1465a9401e		
7	68	6.99588477366255	scus	su3	90d3d2fc7ecc430c9621ece335651a01		
8	167	17.1810699588477	scus				
9	55	5.65843621399177	weu	su4	be1d6d7ac9574cbc9a22cb8ee20f16fc		
10	92	9.46502057613169				10007007	
11	90	9.25925925925926				10007006	
12	57	5.8641975308642					00000000-0000-0000-0000-000000000000

Basket implements the "Apriori" algorithm for item set mining. It extracts all segments whose coverage of the record set is above a threshold (default 5%). You can see that more segments were extracted with similar ones, such as segments 0, 1 or 2, 3.

Both plugins are powerful and easy to use. Their limitation is that they cluster a single record set in an unsupervised manner with no labels. It's unclear whether the extracted patterns characterize the selected record set, anomalous records, or the global record set.

Clustering the difference between two records sets

The diffpatterns() plugin overcomes the limitation of `autocluster` and `basket`. Diffpatterns takes two record sets and extracts the main segments that are different. One set usually contains the anomalous record set being investigated. One is analyzed by `autocluster` and `basket`. The other set contains the reference record set, the baseline.

In the following query, `diffpatterns` finds interesting clusters within the spike's two minutes, which are different from the clusters within the baseline. The baseline window is defined as the eight minutes before 15:00, when the spike started. You extend by a binary column (AB), and specify whether a specific record belongs to the baseline or to the anomalous set. `Diffpatterns` implements a supervised learning algorithm, where the two class labels were generated by the anomalous versus the baseline flag (AB).

[Run the query](#)

Kusto

```
let min_peak_t=datetime(2016-08-23 15:00);
let max_peak_t=datetime(2016-08-23 15:02);
let min_baseline_t=datetime(2016-08-23 14:50);
let max_baseline_t=datetime(2016-08-23 14:58); // Leave a gap between the baseline and the spike to avoid the transition zone.
let splitime=(max_baseline_t+min_peak_t)/2.0;
demo_clustering1
| where (PreciseTimeStamp between(min_baseline_t..max_baseline_t)) or
      (PreciseTimeStamp between(min_peak_t..max_peak_t))
| extend AB=iff(PreciseTimeStamp > splitime, 'Anomaly', 'Baseline')
| evaluate diffpatterns(AB, 'Anomaly', 'Baseline')
```

SegmentId	CountA	CountB	PercentA	PercentB	PercentDiffAB	Region	ScaleUnit	DeploymentId	Tracepoint
0	639	21	65.74	1.7	64.04	eau	su7	b5d1d4df547d4a04ac15885617edba7	
1	167	544	17.18	44.16	26.97	scus			
2	92	356	9.47	28.9	19.43				10007007
3	90	336	9.26	27.27	18.01				10007006
4	82	318	8.44	25.81	17.38	ncus	su1	e24ef436e02b4823ac5d5b1465a9401e	
5	55	252	5.66	20.45	14.8	weu	su4	be1d6d7ac9574cbc9a22cb8ee20f16fc	
6	57	204	5.86	16.56	10.69				

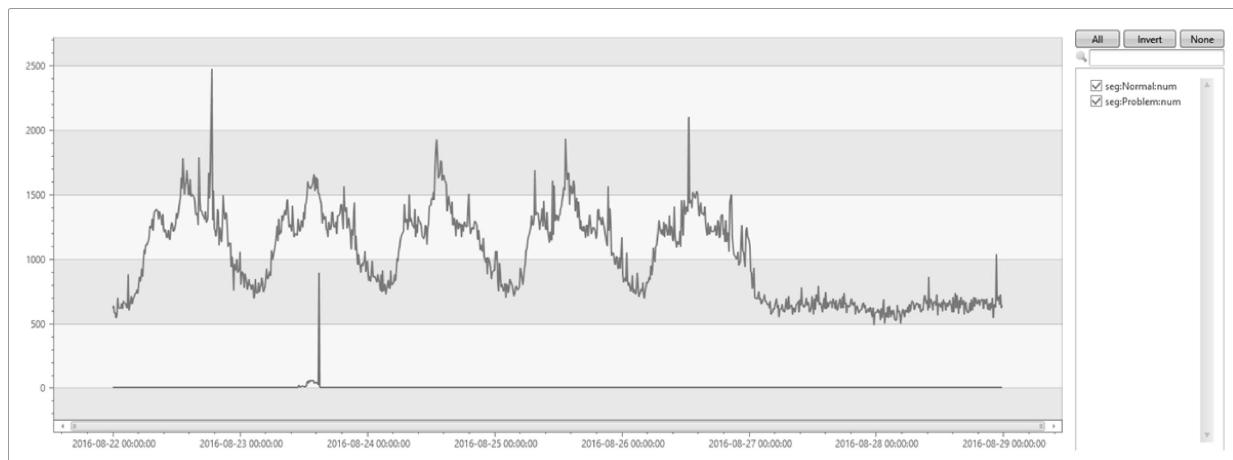
The most dominant segment is the same segment that was extracted by `autocluster`. Its coverage on the two-minute anomalous window is also 65.74%. However, its coverage on the eight-minute baseline window is only 1.7%. The difference is 64.04%. This difference seems to be related to the anomalous spike. To verify this assumption, the following query splits the original chart into the records that belong to this problematic segment, and records from the other segments.

[Run the query](#)

Kusto

```
let min_t = toscalar(demo_clustering1 | summarize min(PreciseTimeStamp));
let max_t = toscalar(demo_clustering1 | summarize max(PreciseTimeStamp));
demo_clustering1
| extend seg = iff(Region == "eau" and ScaleUnit == "su7" and DeploymentId == "b5d1d4df547d4a04ac15885617edba57" and ServiceHost == "e7f60c5d-4944-42b3-922a-92e98a8e7dec", "Problem", "Normal")
```

```
| make-series num=count() on PreciseTimeStamp from min_t to max_t step 10m by seg  
| render timechart
```



This chart allows us to see that the spike on Tuesday afternoon was because of exceptions from this specific segment, discovered by using the `diffpatterns` plugin.

Summary

The Machine Learning plugins are helpful for many scenarios. The `autocluster` and `basket` implement an unsupervised learning algorithm and are easy to use. `Diffpatterns` implements a supervised learning algorithm and, although more complex, it's more powerful for extracting differentiation segments for RCA.

These plugins are used interactively in ad-hoc scenarios and in automatic near real-time monitoring services. Time series anomaly detection is followed by a diagnosis process. The process is highly optimized to meet necessary performance standards.

Feedback

Was this page helpful?

Provide product feedback [↗](#) | Get help at Microsoft Q&A

make-series operator

Article • 01/16/2023

Create series of specified aggregated values along a specified axis.

Syntax

```
T | make-series [MakeSeriesParameters] [Column =] Aggregation [default = DefaultValue] [, ...] on AxisColumn [from start] [to end] step step [by [Column = GroupExpression [, ...]]]
```

Parameters

Name	Type	Required	Description
<i>Column</i>	string		The name for the result column. Defaults to a name derived from the expression.
<i>DefaultValue</i>	scalar		A default value to use instead of absent values. If there's no row with specific values of <i>AxisColumn</i> and <i>GroupExpression</i> , then the corresponding element of the array will be assigned a <i>DefaultValue</i> . Default is 0.
<i>Aggregation</i>	string	✓	A call to an aggregation function, such as <code>count()</code> or <code>avg()</code> , with column names as arguments. See the list of aggregation functions. Only aggregation functions that return numeric results can be used with the <code>make-series</code> operator.
<i>AxisColumn</i>	string	✓	The column by which the series will be ordered. Usually the column values will be of type <code>datetime</code> or <code>timespan</code> but all numeric types are accepted.
<i>start</i>	scalar	✓	The low bound value of the <i>AxisColumn</i> for each of the series to be built. If <i>start</i> is not specified, it will be the first bin, or step, that has data in each series.
<i>end</i>	scalar	✓	The high bound non-inclusive value of the <i>AxisColumn</i> . The last index of the time series is smaller than this value and will be <i>start</i> plus integer multiple of <i>step</i> that is smaller than <i>end</i> . If <i>end</i> is not specified, it will be the upper bound of the last bin, or step, that has data per each series.

Name	Type	Required	Description
<code>step</code>	scalar	✓	The difference, or bin size, between two consecutive elements of the <code>AxisColumn</code> array. For a list of possible time intervals, see <code>timespan</code> .
<code>GroupExpression</code>			An expression over the columns that provides a set of distinct values. Typically it's a column name that already provides a restricted set of values.
<code>MakeSeriesParameters</code>			Zero or more space-separated parameters in the form of <code>Name = Value</code> that control the behavior. See supported make series parameters.

ⓘ Note

The `start`, `end`, and `step` parameters are used to build an array of `AxisColumn` values. The array consists of values between `start` and `end`, with the `step` value representing the difference between one array element to the next. All `Aggregation` values are ordered respectively to this array.

Supported make series parameters

Name	Description
<code>kind</code>	Produces default result when the input of make-series operator is empty. Value: <code>nonempty</code>
<code>hint.shufflekey= <key></code>	The <code>shufflekey</code> query shares the query load on cluster nodes, using a key to partition data. See <code>shuffle</code> query

ⓘ Note

The arrays generated by make-series are limited to 1048576 values (2^{20}). Trying to generate a larger array with make-series would result in either an error or a truncated array.

Alternate Syntax

```
T | make-series [Column =] Aggregation [default = DefaultValue] [, ...] on
AxisColumn in range(start, stop, step) [by [Column =] GroupExpression [, ...]]
```

The generated series from the alternate syntax differs from the main syntax in two aspects:

- The *stop* value is inclusive.
- Binning the index axis is generated with `bin()` and not `bin_at()`, which means that *start* may not be included in the generated series.

It's recommended to use the main syntax of `make-series` and not the alternate syntax.

Returns

The input rows are arranged into groups having the same values of the `by` expressions and the `bin_at(AxisColumn, step, start)` expression. Then the specified aggregation functions are computed over each group, producing a row for each group. The result contains the `by` columns, *AxisColumn* column and also at least one column for each computed aggregate. (Aggregations over multiple columns or non-numeric results aren't supported.)

This intermediate result has as many rows as there are distinct combinations of `by` and `bin_at(AxisColumn, step, start)` values.

Finally the rows from the intermediate result arranged into groups having the same values of the `by` expressions and all aggregated values are arranged into arrays (values of `dynamic` type). For each aggregation, there's one column containing its array with the same name. The last column is an array containing the values of *AxisColumn* binned according to the specified *step*.

ⓘ Note

Although you can provide arbitrary expressions for both the aggregation and grouping expressions, it's more efficient to use simple column names.

List of aggregation functions

Function	Description
<code>avg()</code>	Returns an average value across the group
<code>avgif()</code>	Returns an average with the predicate of the group
<code>count()</code>	Returns a count of the group

Function	Description
countif()	Returns a count with the predicate of the group
dcount()	Returns an approximate distinct count of the group elements
dcountif()	Returns an approximate distinct count with the predicate of the group
max()	Returns the maximum value across the group
maxif()	Returns the maximum value with the predicate of the group
min()	Returns the minimum value across the group
minif()	Returns the minimum value with the predicate of the group
percentile()	Returns the percentile value across the group
take_any()	Returns a random non-empty value for the group
stdev()	Returns the standard deviation across the group
sum()	Returns the sum of the elements within the group
sumif()	Returns the sum of the elements with the predicate of the group
variance()	Returns the variance across the group

List of series analysis functions

Function	Description
series_fir()	Applies Finite Impulse Response \Leftrightarrow filter
series_iir()	Applies Infinite Impulse Response \Leftrightarrow filter
series_fit_line()	Finds a straight line that is the best approximation of the input
series_fit_line_dynamic()	Finds a line that is the best approximation of the input, returning dynamic object
series_fit_2lines()	Finds two lines that are the best approximation of the input
series_fit_2lines_dynamic()	Finds two lines that are the best approximation of the input, returning dynamic object
series_outliers()	Scores anomaly points in a series
series_periods_detect()	Finds the most significant periods that exist in a time series

Function	Description
series_periods_validate()	Checks whether a time series contains periodic patterns of given lengths
series_stats_dynamic()	Return multiple columns with the common statistics (min/max/variance/stdev/average)
series_stats()	Generates a dynamic value with the common statistics (min/max/variance/stdev/average)

For a complete list of series analysis functions, see: Series processing functions

List of series interpolation functions

Function	Description
series_fill_backward()	Performs backward fill interpolation of missing values in a series
series_fill_const()	Replaces missing values in a series with a specified constant value
series_fill_forward()	Performs forward fill interpolation of missing values in a series
series_fill_linear()	Performs linear interpolation of missing values in a series

- Note: Interpolation functions by default assume `null` as a missing value. Therefore specify `default=double(null)` in `make-series` if you intend to use interpolation functions for the series.

Examples

A table that shows arrays of the numbers and average prices of each fruit from each supplier ordered by the timestamp with specified range. There's a row in the output for each distinct combination of fruit and supplier. The output columns show the fruit, supplier, and arrays of: count, average, and the whole timeline (from 2016-01-01 until 2016-01-10). All arrays are sorted by the respective timestamp and all gaps are filled with default values (0 in this example). All other input columns are ignored.

Kusto

```
T | make-series PriceAvg=avg(Price) default=0
on Purchase from datetime(2016-09-10) to datetime(2016-09-13) step 1d by
Supplier, Fruit
```

The diagram illustrates a data processing pipeline. It starts with two input tables on the left, which are then processed by a central transformation step (indicated by a large grey arrow pointing right). The final output is a single table on the right.

Input Table 1:

Supplier	Fruit	Price	Purchase
Aldi	Apple	4	2016-09-10
Costco	Apple	2	2016-09-11
Aldi	Apple	6	2016-09-10
Costco	Snargaluff	100	2016-09-12
Aldi	Apple	7	2016-09-12
Aldi	Snargaluff	400	2016-09-11
Costco	Snargaluff	104	2016-09-12
Aldi	Apple	5	2016-09-12
Aldi	Snargaluff	600	2016-09-11
Costco	Snargaluff	200	2016-09-10

Input Table 2:

Supplier	Fruit	PriceAvg	Purchase
Aldi	Apple	5	2016-09-10
Aldi	Apple	6	2016-09-12
Costco	Apple	2	2016-09-11
Aldi	Snargaluff	500	2016-09-11
Costco	Snargaluff	200	2016-09-10
Costco	Snargaluff	102	2016-09-12

Transformation Step:

The transformation step involves calculating average prices and purchase dates for each fruit across suppliers. The resulting table is:

Suppliers	Fruit	PriceAvg	Purchase
Aldi	Apple	[5,0,6]	[2016-09-10,2016-09-11,2016-09-12]
Costco	Apple	[0,2,0]	[2016-09-10,2016-09-11,2016-09-12]
Aldi	Snargaluff	[0,500,0]	[2016-09-10,2016-09-11,2016-09-12]
Costco	Snargaluff	[200,0,102]	[2016-09-10,2016-09-11,2016-09-12]

Run the query

Kusto

```
let data=datatable(timestamp:datetime, metric: real)
[
    datetime(2016-12-31T06:00), 50,
    datetime(2017-01-01), 4,
    datetime(2017-01-02), 3,
    datetime(2017-01-03), 4,
    datetime(2017-01-03T03:00), 6,
    datetime(2017-01-05), 8,
    datetime(2017-01-05T13:40), 13,
    datetime(2017-01-06), 4,
    datetime(2017-01-07), 3,
    datetime(2017-01-08), 8,
    datetime(2017-01-08T21:00), 8,
    datetime(2017-01-09), 2,
    datetime(2017-01-09T12:00), 11,
    datetime(2017-01-10T05:00), 5,
];
let interval = 1d;
let stime = datetime(2017-01-01);
let etime = datetime(2017-01-10);
data
| make-series avg(metric) on timestamp from stime to etime step interval
```

avg_metric timestamp

avg_metric timestamp

```
[ 4.0, 3.0,      [ "2017-01-01T00:00:00.0000000Z", "2017-01-02T00:00:00.0000000Z", "2017-01-03T00:00:00.0000000Z", "2017-01-04T00:00:00.0000000Z", "2017-01-05T00:00:00.0000000Z", "2017-01-06T00:00:00.0000000Z", "2017-01-07T00:00:00.0000000Z", "2017-01-08T00:00:00.0000000Z", "2017-01-09T00:00:00.0000000Z" ]
```

When the input to `make-series` is empty, the default behavior of `make-series` produces an empty result.

[Run the query](#)

Kusto

```
let data=datatable(timestamp:datetime, metric: real)
[
    datetime(2016-12-31T06:00), 50,
    datetime(2017-01-01), 4,
    datetime(2017-01-02), 3,
    datetime(2017-01-03), 4,
    datetime(2017-01-03T03:00), 6,
    datetime(2017-01-05), 8,
    datetime(2017-01-05T13:40), 13,
    datetime(2017-01-06), 4,
    datetime(2017-01-07), 3,
    datetime(2017-01-08), 8,
    datetime(2017-01-08T21:00), 8,
    datetime(2017-01-09), 2,
    datetime(2017-01-09T12:00), 11,
    datetime(2017-01-10T05:00), 5,
];
let interval = 1d;
let stime = datetime(2017-01-01);
let etime = datetime(2017-01-10);
data
| take 0
| make-series avg(metric) default=1.0 on timestamp from stime to etime step
interval
| count
```

Output

Count

```
0
```

Using `kind=nonempty` in `make-series` will produce a non-empty result of the default values:

Run the query

Kusto

```
let data=datatable(timestamp:datetime, metric: real)
[
    datetime(2016-12-31T06:00), 50,
    datetime(2017-01-01), 4,
    datetime(2017-01-02), 3,
    datetime(2017-01-03), 4,
    datetime(2017-01-03T03:00), 6,
    datetime(2017-01-05), 8,
    datetime(2017-01-05T13:40), 13,
    datetime(2017-01-06), 4,
    datetime(2017-01-07), 3,
    datetime(2017-01-08), 8,
    datetime(2017-01-08T21:00), 8,
    datetime(2017-01-09), 2,
    datetime(2017-01-09T12:00), 11,
    datetime(2017-01-10T05:00), 5,
];
let interval = 1d;
let stime = datetime(2017-01-01);
let etime = datetime(2017-01-10);
data
| take 0
| make-series kind=nonempty avg(metric) default=1.0 on timestamp from stime
to etime step interval
```

Output

avg_metric	timestamp
[[
1.0,	"2017-01-01T00:00:00.0000000Z",
1.0,	"2017-01-02T00:00:00.0000000Z",
1.0,	"2017-01-03T00:00:00.0000000Z",
1.0,	"2017-01-04T00:00:00.0000000Z",
1.0,	"2017-01-05T00:00:00.0000000Z",
1.0,	"2017-01-06T00:00:00.0000000Z",
1.0,	"2017-01-07T00:00:00.0000000Z",
1.0,	"2017-01-08T00:00:00.0000000Z",
1.0	"2017-01-09T00:00:00.0000000Z"
]]

Feedback

Was this page helpful? Yes No

Provide product feedback  | Get help at Microsoft Q&A

series_abs()

Article • 01/31/2023

Calculates the element-wise absolute value of the numeric series input.

Syntax

```
series_abs( series )
```

Parameters

Name	Type	Required	Description
<i>series</i>	dynamic	✓	An array of numeric values over which the absolute value function is applied.

Returns

Dynamic array of calculated absolute value. Any non-numeric element yields a `null` element value.

Example

Run the query

Kusto

```
print arr = dynamic([-6.5,0,8.2])
| extend arr_abs = series_abs(arr)
```

Output

arr	arr_abs
[-6.5,0,8.2]	[6.5,0,8.2]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_acos()

Article • 01/31/2023

Calculates the element-wise arccosine function of the numeric series input.

Syntax

```
series_acos(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the arccosine function is applied.

Returns

Dynamic array of calculated arccosine function values. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
print arr = dynamic([-1,0,1])  
| extend arr_acos = series_acos(arr)
```

Output

arr	arr_acos
[-6.5,0,8.2]	[3.1415926535897931,1.5707963267948966,0.0]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_add()

Article • 01/31/2023

Calculates the element-wise addition of two numeric series inputs.

Syntax

```
series_add(series1, series2)
```

Parameters

Name	Type	Required	Description
series1	dynamic	✓	The numeric arrays to be element-wise added into a dynamic array result.
series2			

Returns

Dynamic array of calculated element-wise add operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
range x from 1 to 3 step 1
| extend y = x * 2
| extend z = y * 2
| project s1 = pack_array(x,y,z), s2 = pack_array(z, y, x)
| extend s1_add_s2 = series_add(s1, s2)
```

Output

s1	s2	s1_add_s2
[1,2,4]	[4,2,1]	[5,4,5]

s1	s2	s1_add_s2
[2,4,8]	[8,4,2]	[10,8,10]
[3,6,12]	[12,6,3]	[15,12,15]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_asin()

Article • 01/31/2023

Calculates the element-wise arcsine function of the numeric series input.

Syntax

```
series_asin(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the arcsine function is applied.

Returns

Dynamic array of calculated arcsine function values. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
print arr = dynamic([-1,0,1])  
| extend arr_asin = series_asin(arr)
```

Output

arr	arr_asin
[-6.5,0,8.2]	[1.5707963267948966,0.0,1.5707963267948966]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_atan()

Article • 01/31/2023

Calculates the element-wise arctangent function of the numeric series input.

Syntax

```
series_atan(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the arctangent function is applied.

Returns

Dynamic array of calculated arctangent function values. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
  
print arr = dynamic([-1,0,1])  
| extend arr_atan = series_atan(arr)
```

Output

arr	arr_atan
[-6.5,0,8.2]	[-0.78539816339744828,0.0,0.78539816339744828]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_ceiling()

Article • 01/31/2023

Calculates the element-wise ceiling function of the numeric series input.

Syntax

```
series_ceiling(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the ceiling function is applied.

Returns

Dynamic array of the calculated ceiling function. Any non-numeric element yields a `null` element value.

Example

Run the query

Kusto

```
print s = dynamic([-1.5,1,2.5])
| extend s_ceiling = series_ceiling(s)
```

Output

s	s_ceiling
[-1.5,1,2.5]	[-1.0,1.0,3.0]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_cos()

Article • 01/31/2023

Calculates the element-wise cosine function of the numeric series input.

Syntax

```
series_cos(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the cosine function is applied.

Returns

Dynamic array of calculated cosine function values. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
print arr = dynamic([-1,0,1])  
| extend arr_cos = series_cos(arr)
```

Output

arr	arr_cos
[-6.5,0,8.2]	[0.54030230586813976,1.0,0.54030230586813976]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_decompose()

Article • 03/12/2023

Applies a decomposition transformation on a series.

Takes an expression containing a series (dynamic numerical array) as input and decomposes it to seasonal, trend, and residual components.

Syntax

```
series_decompose( Series , [ Seasonality , Trend , Test_points , Seasonality_threshold ] )
```

Parameters

Name	Type	Required	Description
<i>Series</i>	dynamic	✓	An array of numeric values, typically the resulting output of make-series or make_list operators.
<i>Seasonality</i>	int		<p>Controls the seasonal analysis. The possible values are:</p> <ul style="list-style-type: none">-1: Autodetect seasonality using series_periods_detect. This is the default value.- Period: A positive integer specifying the expected period in number of bins. For example, if the series is in 1 - h bins, a weekly period is 168 bins.0: No seasonality, so skip extracting this component.
<i>Trend</i>	string		<p>Controls the trend analysis. The possible values are:</p> <ul style="list-style-type: none">avg: Define trend component as average(x). This is the default.linefit: Extract trend component using linear regression.none: No trend, so skip extracting this component.

Name	Type	Required	Description
<i>Test_points</i>	int		A positive integer specifying the number of points at the end of the series to exclude from the learning, or regression, process. This parameter should be set for forecasting purposes. The default value is 0.
<i>Seasonality_threshold</i>	real		The threshold for seasonality score when <i>Seasonality</i> is set to autodetect. The default score threshold is 0.6.
For more information, see series_periods_detect .			

Returns

The function returns the following respective series:

- `baseline`: the predicted value of the series (sum of seasonal and trend components, see below).
- `seasonal`: the series of the seasonal component:
 - if the period isn't detected or is explicitly set to 0: constant 0.
 - if detected or set to positive integer: median of the series points in the same phase
- `trend`: the series of the trend component.
- `residual`: the series of the residual component (that is, $x - \text{baseline}$).

ⓘ Note

- Component execution order:
 1. Extract the seasonal series
 2. Subtract it from x , generating the deseasonal series
 3. Extract the trend component from the deseasonal series
 4. Create the baseline = seasonal + trend
 5. Create the residual = $x - \text{baseline}$
- Either seasonality and, or trend should be enabled. Otherwise, the function is redundant, and just returns $\text{baseline} = 0$ and $\text{residual} = x$.

[More about series decomposition](#)

This method is usually applied to time series of metrics expected to manifest periodic and/or trend behavior. You can use the method to forecast future metric values and/or detect anomalous values. The implicit assumption of this regression process is that apart from seasonal and trend behavior, the time series is stochastic and randomly distributed. Forecast future metric values from the seasonal and trend components while ignoring the residual part. Detect anomalous values based on outlier detection only on the residual part only. Further details can be found in the Time Series Decomposition chapter ↗.

Examples

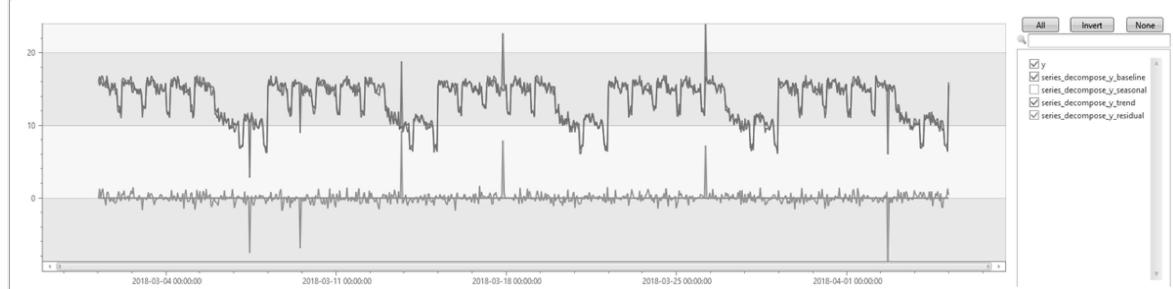
Weekly seasonality

In the following example, we generate a series with weekly seasonality and without trend, we then add some outliers to it. `series_decompose` finds and automatically detects the seasonality, and generates a baseline that is almost identical to the seasonal component. The outliers we added can be clearly seen in the residuals component.

Run the query

```
Kusto

let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 10.0, 15.0) - (((t%24)/10)*
((t%24)/10)) // generate a series with weekly seasonality
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose(y)
| render timechart
```



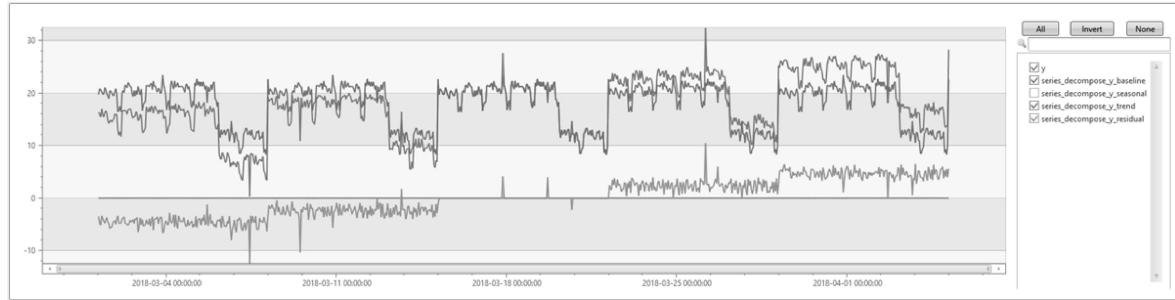
Weekly seasonality with trend

In this example, we add a trend to the series from the previous example. First, we run `series_decompose` with the default parameters. The trend `avg` default value only takes the average and doesn't compute the trend. The generated baseline doesn't contain the trend. When observing the trend in the residuals, it becomes apparent that this example is less accurate than the previous example.

[Run the query](#)

```
Kusto

let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 5.0, 15.0) - (((t%24)/10)*
((t%24)/10)) + t/72.0 // generate a series with weekly seasonality and
ongoing trend
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose(y)
| render timechart
```



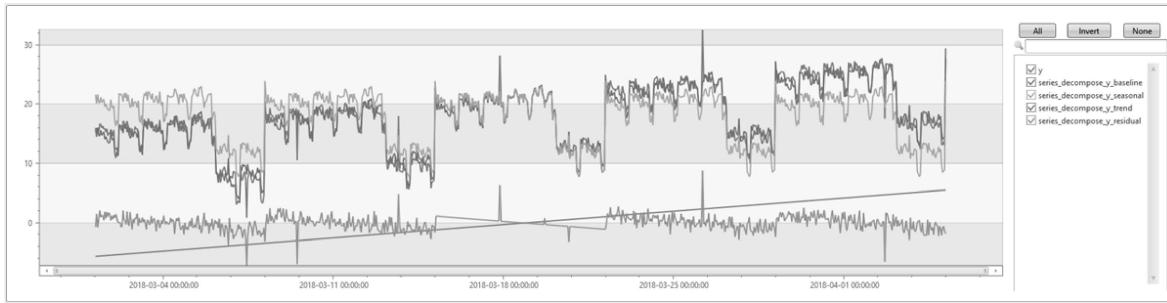
Next, we rerun the same example. Since we're expecting a trend in the series, we specify `linefit` in the trend parameter. We can see that the positive trend is detected and the baseline is much closer to the input series. The residuals are close to zero, and only the outliers stand out. We can see all the components on the series in the chart.

[Run the query](#)

```
Kusto

let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 5.0, 15.0) - (((t%24)/10)*
```

```
((t%24)/10)) + t/72.0 // generate a series with weekly seasonality and
ongoing trend
| extend y=iiff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iiff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose(y, -1, 'linefit')
| render timechart
```



See also

- Visualize results with an anomalychart

Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_decompose_anomalies()

Article • 03/12/2023

Anomaly Detection is based on series decomposition. For more information, see `series_decompose()`.

The function takes an expression containing a series (dynamic numerical array) as input, and extracts anomalous points with scores.

Syntax

```
series_decompose_anomalies (Series, [Threshold, Seasonality, Trend, Test_points,  
AD_method, Seasonality_threshold])
```

Parameters

Name	Type	Required	Description
<i>Series</i>	dynamic	✓	An array of numeric values, typically the resulting output of <code>make-series</code> or <code>make_list</code> operators.
<i>Threshold</i>	real		The anomaly threshold. The default is 1.5, k value, for detecting mild or stronger anomalies.
<i>Seasonality</i>	int		Controls the seasonal analysis. The possible values are: <ul style="list-style-type: none">- <code>-1</code>: Autodetect seasonality using <code>series_periods_detect</code>. This is the default value.- <code>Period</code>: A positive integer specifying the expected period in number of bins. For example, if the series is in <code>1 - h</code> bins, a weekly period is 168 bins.- <code>0</code>: No seasonality, so skip extracting this component.
<i>Trend</i>	string		Controls the trend analysis. The possible values are: <ul style="list-style-type: none">- <code>avg</code>: Define trend component as <code>average(x)</code>. This is the default.- <code>linefit</code>: Extract trend component using linear regression.- <code>none</code>: No trend, so skip extracting this component.

Name	Type	Required	Description
<i>Test_points</i>	int		A positive integer specifying the number of points at the end of the series to exclude from the learning, or regression, process. This parameter should be set for forecasting purposes. The default value is 0.
<i>AD_method</i>	string		<p>Controls the anomaly detection method on the residual time series, containing one of the following values:</p> <ul style="list-style-type: none"> - <code>ctukey</code>: Tukey's fence test ↗ with custom 10th-90th percentile range. This is the default. - <code>tukey</code>: Tukey's fence test ↗ with standard 25th-75th percentile range. <p>For more information on residual time series, see <code>series_outliers</code>.</p>
<i>Seasonality_threshold</i>	real		<p>The threshold for seasonality score when <code>Seasonality</code> is set to autodetect. The default score threshold is 0.6.</p> <p>For more information, see <code>series_periods_detect</code>.</p>

Returns

The function returns the following respective series:

- `ad_flag`: A ternary series containing (+1, -1, 0) marking up/down/no anomaly respectively
- `ad_score`: Anomaly score
- `baseline`: The predicted value of the series, according to the decomposition

The algorithm

This function follows these steps:

1. Calls `series_decompose()` with the respective parameters, to create the baseline and residuals series.
2. Calculates `ad_score` series by applying `series_outliers()` with the chosen anomaly detection method on the residuals series.

3. Calculates the ad_flag series by applying the threshold on the ad_score to mark up/down/no anomaly respectively.

Examples

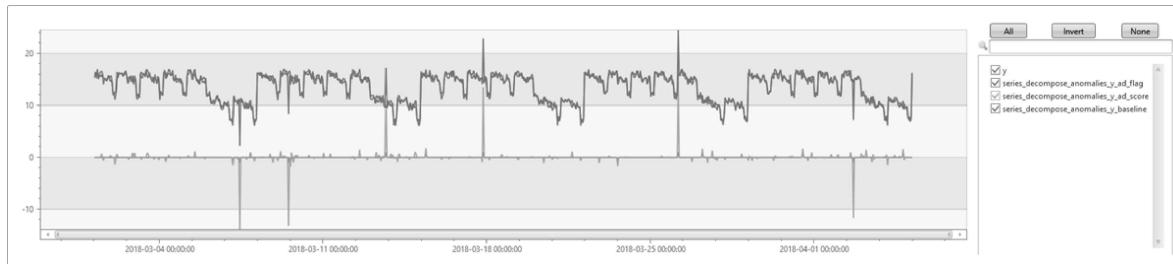
Detect anomalies in weekly seasonality

In the following example, generate a series with weekly seasonality, and then add some outliers to it. `series_decompose_anomalies` autodetects the seasonality and generates a baseline that captures the repetitive pattern. The outliers you added can be clearly spotted in the ad_score component.

[Run the query](#)

Kusto

```
let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 10.0, 15.0) - (((t%24)/10)*
((t%24)/10)) // generate a series with weekly seasonality
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose_anomalies(y)
| render timechart
```



Detect anomalies in weekly seasonality with trend

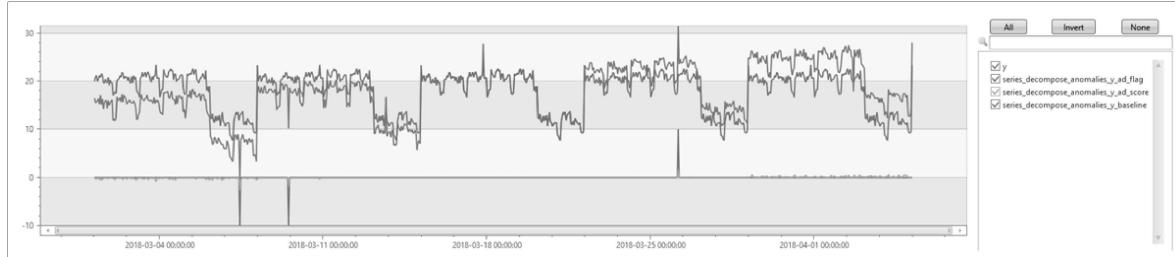
In this example, add a trend to the series from the previous example. First, run `series_decompose_anomalies` with the default parameters in which the trend `avg` default value only takes the average and doesn't compute the trend. The generated baseline doesn't contain the trend and is less exact, compared to the previous example.

Consequently, some of the outliers you inserted in the data aren't detected because of the higher variance.

Run the query

Kusto

```
let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 5.0, 15.0) - (((t%24)/10)*
((t%24)/10)) + t/72.0 // generate a series with weekly seasonality and
ongoing trend
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose_anomalies(y)
| extend series_decompose_anomalies_y_ad_flag =
series_multiply(10, series_decompose_anomalies_y_ad_flag) // multiply by 10
for visualization purposes
| render timechart
```



Next, run the same example, but since you're expecting a trend in the series, specify `linefit` in the trend parameter. You can see that the baseline is much closer to the input series. All the inserted outliers are detected, and also some false positives. See the next example on tweaking the threshold.

Run the query

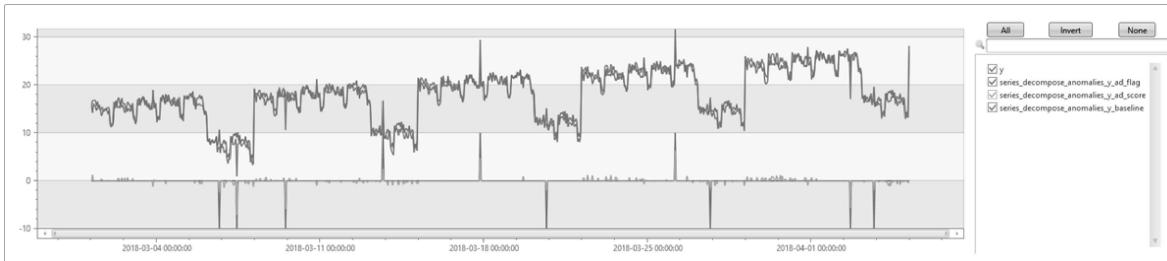
Kusto

```
let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 5.0, 15.0) - (((t%24)/10)*
((t%24)/10)) + t/72.0 // generate a series with weekly seasonality and
ongoing trend
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
```

```

outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose_anomalies(y, 1.5, -1, 'linefit')
| extend series_decompose_anomalies_y_ad_flag =
series_multiply(10, series_decompose_anomalies_y_ad_flag) // multiply by 10
for visualization purposes
| render timechart

```



Tweak the anomaly detection threshold

A few noisy points were detected as anomalies in the previous example. Now increase the anomaly detection threshold from a default of 1.5 to 2.5. Use this interpercentile range, so that only stronger anomalies are detected. Now, only the outliers you inserted in the data, will be detected.

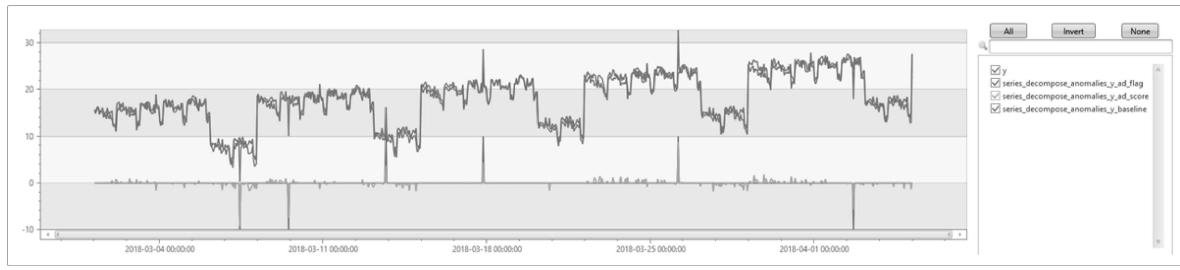
[Run the query](#)

Kusto

```

let ts=range t from 1 to 24*7*5 step 1
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 5.0, 15.0) - (((t%24)/10)*
((t%24)/10)) + t/72.0 // generate a series with weekly seasonality and
ongoing trend
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
| summarize Timestamp=make_list(Timestamp, 10000),y=make_list(y, 10000);
ts
| extend series_decompose_anomalies(y, 2.5, -1, 'linefit')
| extend series_decompose_anomalies_y_ad_flag =
series_multiply(10, series_decompose_anomalies_y_ad_flag) // multiply by 10
for visualization purposes
| render timechart

```



Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_decompose_forecast()

Article • 03/12/2023

Forecast based on series decomposition.

Takes an expression containing a series (dynamic numerical array) as input, and predicts the values of the last trailing points. For more information, see [series_decompose](#).

Syntax

```
series_decompose_forecast( Series , Points , [ Seasonality , Trend ,  
Seasonality_threshold ] )
```

Parameters

Name	Type	Required	Description
<i>Series</i>	dynamic	✓	An array of numeric values, typically the resulting output of make-series or make_list operators.
<i>Points</i>	int	✓	Specifies the number of points at the end of the series to predict, or forecast. These points are excluded from the learning, or regression, process.
<i>Seasonality</i>	int		Controls the seasonal analysis. The possible values are: <ul style="list-style-type: none">- <code>-1</code>: Autodetect seasonality using <code>series_periods_detect</code>. This is the default value.- <code>Period</code>: A positive integer specifying the expected period in number of bins. For example, if the series is in <code>1 - h</code> bins, a weekly period is 168 bins.- <code>0</code>: No seasonality, so skip extracting this component.
<i>Trend</i>	string		Controls the trend analysis. The possible values are: <ul style="list-style-type: none">- <code>avg</code>: Define trend component as <code>average(x)</code>. This is the default.- <code>linefit</code>: Extract trend component using linear regression.- <code>none</code>: No trend, so skip extracting this component.

Name	Type	Required	Description
<i>Seasonality_threshold</i>	real		The threshold for seasonality score when <i>Seasonality</i> is set to autodetect. The default score threshold is 0.6. For more information, see series_periods_detect .

Returns

A dynamic array with the forecasted series.

ⓘ Note

- The dynamic array of the original input series should include a number of *points* slots to be forecasted. The forecast is typically done by using `make-series` and specifying the end time in the range that includes the timeframe to forecast.
- Either seasonality or trend should be enabled, otherwise the function is redundant, and just returns a series filled with zeroes.

Example

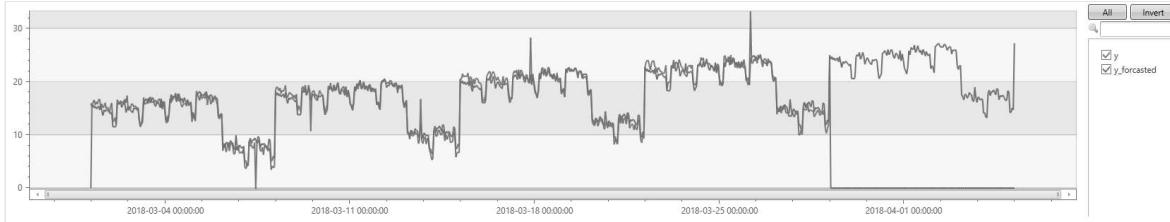
In the following example, we generate a series of four weeks in an hourly grain, with weekly seasonality and a small upward trend. We then use `make-series` and add another empty week to the series. `series_decompose_forecast` is called with a week (24*7 points), and it automatically detects the seasonality and trend, and generates a forecast of the entire five-week period.

[Run the query](#)

Kusto

```
let ts=range t from 1 to 24*7*4 step 1 // generate 4 weeks of hourly data
| extend Timestamp = datetime(2018-03-01 05:00) + 1h * t
| extend y = 2*rand() + iff((t/24)%7>=5, 5.0, 15.0) - (((t%24)/10)*
((t%24)/10)) + t/72.0 // generate a series with weekly seasonality and
ongoing trend
| extend y=iff(t==150 or t==200 or t==780, y-8.0, y) // add some dip
outliers
| extend y=iff(t==300 or t==400 or t==600, y+8.0, y) // add some spike
outliers
```

```
| make-series y=max(y) on Timestamp from datetime(2018-03-01 05:00) to  
datetime(2018-03-01 05:00)+24*7*5h step 1h; // create a time series of 5  
weeks (last week is empty)  
ts  
| extend y_forcasted = series_decompose_forecast(y, 24*7) // forecast a  
week forward  
| render timechart
```



Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_divide()

Article • 01/31/2023

Calculates the element-wise division of two numeric series inputs.

Syntax

```
series_divide(series1, series2)
```

Parameters

Name	Type	Required	Description
<i>series1</i>	dynamic	✓	The numeric arrays over which to calculate the element-wise division. The first array is to be divided by the second.
<i>series2</i>			

Returns

Dynamic array of calculated element-wise divide operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Note: the result series is of double type, even if the inputs are integers. Division by zero follows the double division by zero (e.g. 2/0 yields double(+inf)).

Example

Run the query

Kusto

```
range x from 1 to 3 step 1
| extend y = x * 2
| extend z = y * 2
| project s1 = pack_array(x,y,z), s2 = pack_array(z, y, x)
| extend s1_divide_s2 = series_divide(s1, s2)
```

Output

s1	s2	s1_divide_s2
[1,2,4]	[4,2,1]	[0.25,1.0,4.0]
[2,4,8]	[8,4,2]	[0.25,1.0,4.0]
[3,6,12]	[12,6,3]	[0.25,1.0,4.0]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_dot_product()

Article • 03/12/2023

Calculates the dot product of two numeric series.

The function `series_dot_product()` takes two numeric series as input, and calculates their dot product[↗].

Syntax

```
series_dot_product(series1, series2)
```

Alternate syntax

```
series_dot_product(series, numeric)
```

```
series_dot_product(numeric, series)
```

ⓘ Note

The alternate syntax shows that one of the two function arguments can be a numerical scalar.

This numerical scalar will be broadcasted to a vector whose length equals the length of the corresponding numeric series.

For example, `series_dot_product([1, 2, 3], 10)` will be treated as
`series_dot_product([1, 2, 3], [10, 10, 10]).`

Parameters

Name	Type	Required	Description
<code>series1</code> , <code>series2</code>	dynamic	✓	Input arrays with numeric data, to be element-wise multiplied and then summed into a value of type <code>real</code> .

Returns

Returns a value of type `real` whose value is the sum over the product of each element of *series1* with the corresponding element of *series2*. In case both series length isn't equal, the longer series will be truncated to the length of the shorter one. Any non-numeric element of the input series will be ignored.

ⓘ Note

If one or both input arrays are empty, the result will be `null`.

Example

Run the query

Kusto

```
range x from 1 to 3 step 1
| extend y = x * 2
| extend z = y * 2
| project s1 = pack_array(x,y,z), s2 = pack_array(z, y, x)
| extend s1_dot_product_s2 = series_dot_product(s1, s2)
```

s1	s2	s1_dot_product_s2
[1,2,4]	[4,2,1]	12
[2,4,8]	[8,4,2]	48
[3,6,12]	[12,6,3]	108

Run the query

Kusto

```
range x from 1 to 3 step 1
| extend y = x * 2
| extend z = y * 2
| project s1 = pack_array(x,y,z), s2 = x
| extend s1_dot_product_s2 = series_dot_product(s1, s2)
```

s1	s2	s1_dot_product_s2
[1,2,4]	1	7

s1	s2	s1_dot_product_s2
[2,4,8]	2	28
[3,6,12]	3	63

Feedback

Was this page helpful?



Yes



No

Provide product feedback | Get help at Microsoft Q&A

series_equals()

Article • 01/31/2023

Calculates the element-wise equals (==) logic operation of two numeric series inputs.

Syntax

```
series_equals (series1, series2)
```

Parameters

Name	Type	Required	Description
series1, series2	dynamic	✓	The numeric arrays to be element-wise compared.

Returns

Dynamic array of booleans containing the calculated element-wise equal logic operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
print s1 = dynamic([1,2,4]), s2 = dynamic([4,2,1])
| extend s1_equals_s2 = series_equals(s1, s2)
```

Output

s1	s2	s1_equals_s2
[1,2,4]	[4,2,1]	[false,true,false]

See also

For entire series statistics comparisons, see:

- `series_stats()`
 - `series_stats_dynamic()`
-

Feedback

Was this page helpful?



Yes



No

Provide product feedback | Get help at Microsoft Q&A

series_exp()

Article • 02/01/2023

Calculates the element-wise base-e exponential function (e^x) of the numeric series input.

Syntax

```
series_exp(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values whose elements are applied as the exponent in the exponential function.

Returns

Dynamic array of calculated exponential function. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
print s = dynamic([1,2,3])  
| extend s_exp = series_exp(s)
```

Output

s	s_exp
[1,2,3]	[2.7182818284590451,7.38905609893065,20.085536923187668]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_fft()

Article • 02/01/2023

Applies the Fast Fourier Transform (FFT) on a series.

The `series_fft()` function takes a series of complex numbers in the time/spatial domain and transforms it to the frequency domain using the Fast Fourier Transform[↗]. The transformed complex series represents the magnitude and phase of the frequencies appearing in the original series. Use the complementary function `series_ifft` to transform from the frequency domain back to the time/spatial domain.

Syntax

```
series_fft(x_real [, x_imaginary])
```

Parameters

Name	Type	Required	Description
<code>x_real</code>	dynamic	✓	A numeric array representing the real component of the series to transform.
<code>x_imaginary</code>	dynamic		A similar array representing the imaginary component of the series. This parameter should only be specified if the input series contains complex numbers.

Returns

The function returns the complex inverse fft in two series. The first series for the real component and the second one for the imaginary component.

Example

- Generate a complex series, where the real and imaginary components are pure sine waves in different frequencies. Use FFT to transform it to the frequency domain:

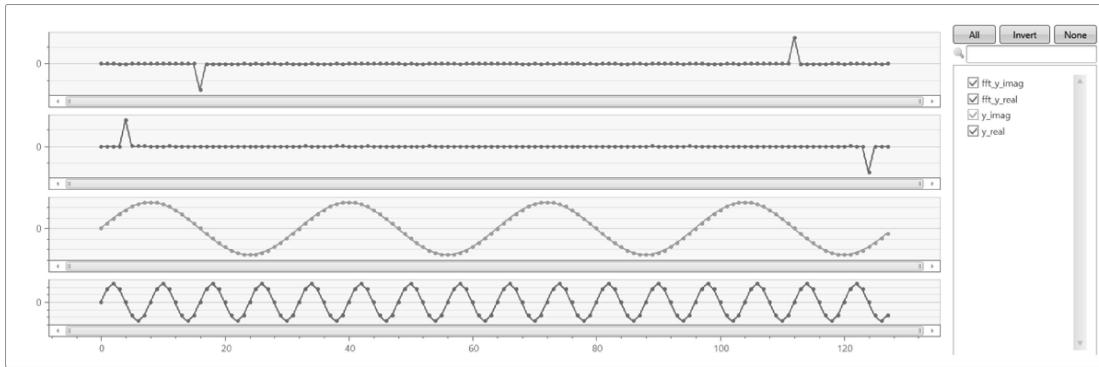
[Run the query](#)

```

let sinewave=(x:double, period:double, gain:double=1.0,
phase:double=0.0)
{
    gain*sin(2*pi()/period*(x+phase))
}
;
let n=128;      //  signal length
range x from 0 to n-1 step 1 | extend yr=sinewave(x, 8), yi=sinewave(x,
32)
| summarize x=make_list(x), y_real=make_list(yr), y_imag=make_list(yi)
| extend (fft_y_real, fft_y_imag) = series_fft(y_real, y_imag)
| render linechart with(ysplit=panels)

```

This query returns *fft_y_real* and *fft_y_imag*:



- Transform a series to the frequency domain, and then apply the inverse transform to get back the original series:

Run the query

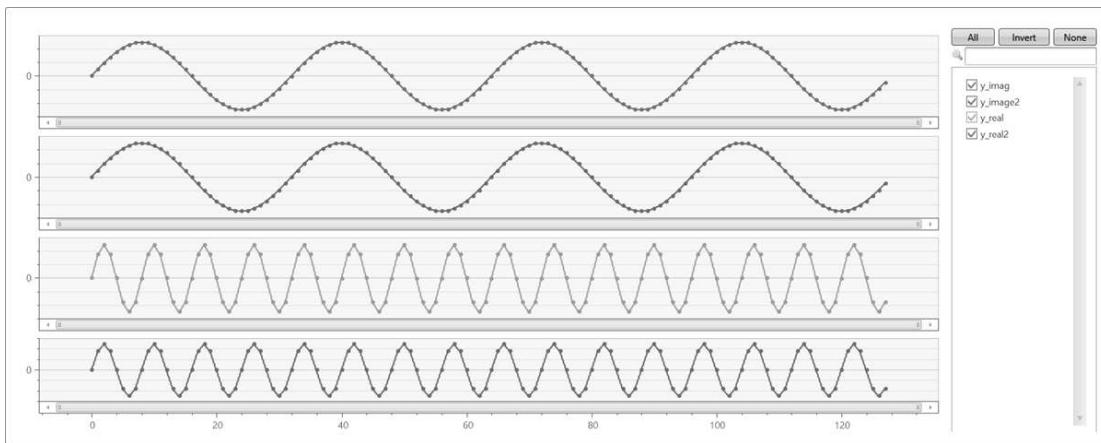
Kusto

```

let sinewave=(x:double, period:double, gain:double=1.0,
phase:double=0.0)
{
    gain*sin(2*pi()/period*(x+phase))
}
;
let n=128;      //  signal length
range x from 0 to n-1 step 1 | extend yr=sinewave(x, 8), yi=sinewave(x,
32)
| summarize x=make_list(x), y_real=make_list(yr), y_imag=make_list(yi)
| extend (fft_y_real, fft_y_imag) = series_fft(y_real, y_imag)
| extend (y_real2, y_image2) = series_ifft(fft_y_real, fft_y_imag)
| project-away fft_y_real, fft_y_imag //  too many series for
linechart with panels
| render linechart with(ysplit=panels)

```

This query returns *y_real2* and **y_imag2*, which are the same as *y_real* and *y_imag*:



Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_fill_backward()

Article • 02/01/2023

Performs a backward fill interpolation of missing values in a series.

An expression containing dynamic numerical array is the input. The function replaces all instances of `missing_value_placeholder` with the nearest value from its right side (other than `missing_value_placeholder`), and returns the resulting array. The rightmost instances of `missing_value_placeholder` are preserved.

Syntax

```
series_fill_backward(series [, missing_value_placeholder])
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>missing_value_placeholder</code>	scalar		Specifies a placeholder for missing values. The default value is <code>double(null)</code> . The value can be of any type that will be converted to actual element types. <code>double(null)</code> , <code>long(null)</code> and <code>int(null)</code> have the same meaning.

ⓘ Note

- If you create `series` using the make-series operator, specify `null` as the default value to use interpolation functions like `series_fill_backward()` afterwards. See explanation.
- If `missing_value_placeholder` is `double(null)`, or omitted, then a result may contain `null` values. To fill these `null` values, use other interpolation functions. Only `series_outliers()` supports `null` values in input arrays.
- `series_fill_backward()` preserves the original type of the array elements.

Returns

series with all instances of *missing_value_placeholder* filled backwards.

Example

[Run the query](#)

Kusto

```
let data = datatable(arr: dynamic)
[
    dynamic([111, null, 36, 41, null, null, 16, 61, 33, null, null])
];
data
| project
    arr,
    fill_backward = series_fill_backward(arr)
```

Output

arr	fill_backward
[111,null,36,41,null,null,16,61,33,null,null]	[111,36,36,41,16,16,16,61,33,null,null]

💡 Tip

Use `series_fill_forward` or `series-fill-const` to complete interpolation of the above array.

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback ↗ | Get help at Microsoft Q&A

series_fill_const()

Article • 02/01/2023

Replaces missing values in a series with a specified constant value.

Takes an expression containing dynamic numerical array as input, replaces all instances of `missing_value_placeholder` with the specified `constant_value` and returns the resulting array.

Syntax

```
series_fill_const(series, constant_value, [missing_value_placeholder])
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>constant_value</code>	scalar	✓	The value used to replace the missing values.
<code>missing_value_placeholder</code>	scalar		Specifies a placeholder for missing values. The default value is <code>double(null)</code> . The value can be of any type that will be converted to actual element types. <code>double(null)</code> , <code>long(null)</code> and <code>int(null)</code> have the same meaning.

Returns

`series` with all instances of `missing_value_placeholder` replaced with `constant_value`.

① Note

- If you create `series` using the `make-series` operator, specify `null` as the default value to use interpolation functions like `series_fill_const()` afterwards. See [explanation](#).
- If `missing_value_placeholder` is `double(null)`, or omitted, then a result may contain `null` values. To fill these `null` values, use other interpolation functions. Only `series_outliers()` supports `null` values in input arrays.
- `series_fill_const()` preserves the original type of the array elements.

Example

Run the query

Kusto

```
let data = datatable(arr: dynamic)
[
    dynamic([111, null, 36, 41, 23, null, 16, 61, 33, null, null])
];
data
| project
    arr,
    fill_const1 = series_fill_const(arr, 0.0),
    fill_const2 = series_fill_const(arr, -1)
```

Output

arr	fill_const1	fill_const2
[111,null,36,41,23,null,16,61,33,null,null]	[111,0.0,36,41,23,0.0,16,61,33,0.0,0.0]	[111,-1,36,41,23,-1,16,61,33,-1,-1]

Feedback

Was this page helpful?  Yes  No

Provide product feedback  | Get help at Microsoft Q&A

series_fill_forward()

Article • 02/01/2023

Performs a forward fill interpolation of missing values in a series.

An expression containing dynamic numerical array is the input. The function replaces all instances of `missing_value_placeholder` with the nearest value from its left side other than `missing_value_placeholder`, and returns the resulting array. The leftmost instances of `missing_value_placeholder` are preserved.

Syntax

```
series_fill_forward(series, [missing_value_placeholder])
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>missing_value_placeholder</code>	scalar		Specifies a placeholder for missing values. The default value is <code>double(null)</code> . The value can be of any type that will be converted to actual element types. <code>double(null)</code> , <code>long(null)</code> and <code>int(null)</code> have the same meaning.

Returns

`series` with all instances of `missing_value_placeholder` filled forwards.

ⓘ Note

- If you create `series` using the `make-series` operator, specify `null` as the default value to use interpolation functions like `series_fill_forward()` afterwards.
See [explanation](#).
- If `missing_value_placeholder` is `double(null)`, or omitted, then a result may contain `null` values. To fill these `null` values, use other interpolation functions.
`Only series_outliers()` supports `null` values in input arrays.
- `series_fill_forward()` preserves the original type of the array elements.

Example

Run the query

Kusto

```
let data = datatable(arr: dynamic)
[
    dynamic([null, null, 36, 41, null, null, 16, 61, 33, null, null])
];
data
| project
    arr,
    fill_forward = series_fill_forward(arr)
```

Output

arr	fill_forward
[null,null,36,41,null,null,16,61,33,null,null]	[null,null,36,41,41,41,16,61,33,33,33]

Use series_fill_backward or series-fill-const to complete interpolation of the above array.

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_fill_linear()

Article • 02/01/2023

Linearly interpolates missing values in a series.

Takes an expression containing dynamic numerical array as input, does linear interpolation for all instances of `missing_value_placeholder`, and returns the resulting array. If the beginning and end of the array contain `missing_value_placeholder`, then it will be replaced with the nearest value other than `missing_value_placeholder`. This feature can be turned off. If the whole array consists of the `missing_value_placeholder`, the array will be filled with `constant_value`, or 0 if not specified.

Syntax

```
series_fill_linear(series, [ missing_value_placeholder [, fill_edges [, constant_value ]]] )
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>missing_value_placeholder</code>	scalar		Specifies a placeholder for missing values. The default value is <code>double(null)</code> . The value can be of any type that will be converted to actual element types. <code>double(null)</code> , <code>long(null)</code> and <code>int(null)</code> have the same meaning.
<code>fill_edges</code>	bool		Indicates whether <code>missing_value_placeholder</code> at the start and end of the array should be replaced with nearest value. <code>true</code> by default. If set to <code>false</code> , then <code>missing_value_placeholder</code> at the start and end of the array will be preserved.
<code>constant_value</code>	scalar		Relevant only for arrays that entirely consist of <code>null</code> values. This parameter specifies a constant value with which to fill the series. Default value is 0. Setting this parameter to <code>double(null)</code> will preserve the <code>null</code> values.

Returns

A series linear interpolation of `series` using the specified parameters. If `series` contains only `int` or `long` elements, then the linear interpolation will return rounded interpolated values rather than exact ones.

ⓘ Note

- If you create `series` using the make-series operator, specify `null` as the default value to use interpolation functions like `series_fill_linear()` afterwards. See [explanation](#).
- If `missing_value_placeholder` is `double(null)`, or omitted, then a result may contain `null` values. To fill these `null` values, use other interpolation functions. Only `series_outliers()` supports `null` values in input arrays.
- `series_fill_linear()` preserves the original type of the array elements.

Example

Run the query

Kusto

```
let data = datatable(arr: dynamic)
[
    dynamic([null, 111.0, null, 36.0, 41.0, null, null, 16.0, 61.0, 33.0, null, null]), // Array of double
    dynamic([null, 111, null, 36, 41, null, null, 16, 61, 33, null, null]), // Similar array of int
    dynamic([null, null, null, null]) // Array with missing values only
];
data
| project
    arr,
    without_args = series_fill_linear(arr),
    with_edges = series_fill_linear(arr, double(null), true),
    wo_edges = series_fill_linear(arr, double(null), false),
    with_const = series_fill_linear(arr, double(null), true, 3.14159)
```

Output

arr	without_args	with_edges
[null,111.0,null,36.0,41.0,null,null,16.0,61.0,33.0,null,null]	[111.0,111.0,73.5,36.0,41.0,32.667,24.333,16.0,61.0,33.0,33.0,33.0]	[111.0,111.0,73.5,36.0,41.0,32.667,24.333,1]
[null,111,null,36,41,null,null,16,61,33,null,null]	[111,111,73,36,41,32,24,16,61,33,33,33]	[111,111,73,36,41,32,24,16,61,33,33,33]
[null,null,null,null]	[0.0.0.0.0.0.0]	[0.0.0.0.0.0.0]

Feedback

Was this page helpful?

Provide product feedback [↗](#) | Get help at Microsoft Q&A

series_fir()

Article • 02/01/2023

Applies a Finite Impulse Response (FIR) filter on a series.

The function takes an expression containing a dynamic numerical array as input and applies a Finite Impulse Response  filter. By specifying the `filter` coefficients, it can be used for calculating a moving average, smoothing, change-detection, and many more use cases. The function takes the column containing the dynamic array and a static dynamic array of the filter's coefficients as input, and applies the filter on the column. It outputs a new dynamic array column, containing the filtered output.

Syntax

```
series_fir(series, filter [, normalize[, center]])
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>filter</code>	dynamic	✓	An array of numeric values containing the coefficients of the filter.
<code>normalize</code>	bool		Indicates whether the filter should be normalized. That is, divided by the sum of the coefficients. If <code>filter</code> contains negative values, then <code>normalize</code> must be specified as <code>false</code> , otherwise result will be <code>null</code> . If not specified, then a default value of <code>true</code> is assumed, depending on the presence of negative values in the <code>filter</code> . If <code>filter</code> contains at least one negative value, then <code>normalize</code> is assumed to be <code>false</code> .
<code>center</code>	bool		Indicates whether the filter is applied symmetrically on a time window before and after the current point, or on a time window from the current point backwards. By default, <code>center</code> is <code>false</code> , which fits the scenario of streaming data so that we can only apply the filter on the current and older points. However, for ad-hoc processing you can set it to <code>true</code> , keeping it synchronized with the time series. See examples below. This parameter controls the filter's group delay  .

 Tip

Normalization is a convenient way to make sure that the sum of the coefficients is 1. When *normalized* is `true`, the filter doesn't amplify or attenuate the series. For example, the moving average of four bins could be specified by `filter=[1,1,1,1]` and `normalized=true`, which is simpler than typing `[0.25,0.25,0.25,0.25]`.

Returns

A new dynamic array column containing the filtered output.

Examples

- Calculate a moving average of five points by setting `filter=[1,1,1,1,1]` and `normalize=true` (default). Note the effect of `center=false` (default) vs. `true`:

Run the query

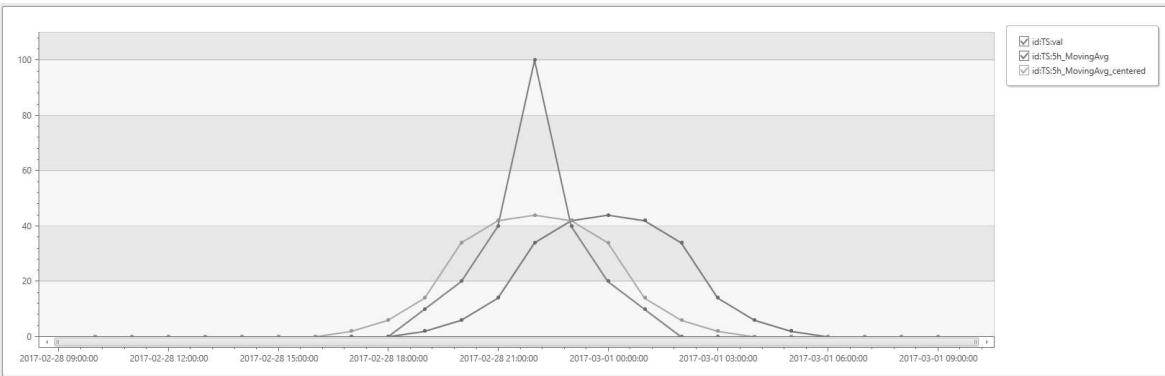
Kusto

```
range t from bin(now(), 1h) - 23h to bin(now(), 1h) step 1h
| summarize t=make_list(t)
| project
    id='TS',
    val=dynamic([0, 0, 0, 0, 0, 0, 0, 0, 0, 10, 20, 40, 100, 40, 20, 10, 0,
0, 0, 0, 0, 0, 0]),
    t
| extend
    5h_MovingAvg=series_fir(val, dynamic([1, 1, 1, 1, 1])),
    5h_MovingAvg_centered=series_fir(val, dynamic([1, 1, 1, 1, 1]), true,
true)
| render timechart
```

This query returns:

`5h_MovingAvg`: Five points moving average filter. The spike is smoothed and its peak shifted by $(5-1)/2 = 2h$.

`5h_MovingAvg_centered`: Same, but by setting `center=true`, the peak stays in its original location.

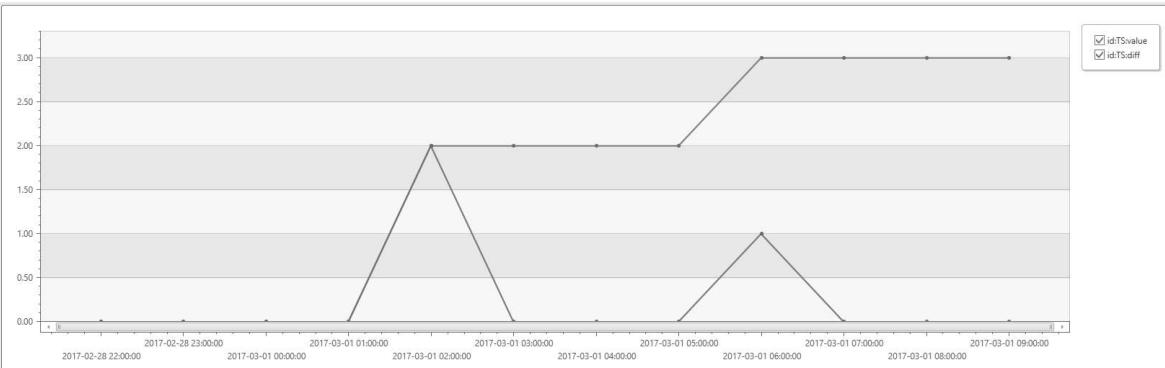


- To calculate the difference between a point and its preceding one, set `filter=[1,-1]`.

Run the query

Kusto

```
range t from bin(now(), 1h) - 11h to bin(now(), 1h) step 1h
| summarize t=make_list(t)
| project id='TS', t, value=dynamic([0, 0, 0, 0, 2, 2, 2, 2, 3, 3, 3])
| extend diff=series_fir(value, dynamic([1, -1]), false, false)
| render timechart
```



Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_fit_2lines()

Article • 02/01/2023

Applies a two segmented linear regression on a series, returning multiple columns.

Takes an expression containing dynamic numerical array as input and applies a two segmented linear regression ↗ in order to identify and quantify a trend change in a series. The function iterates on the series indexes. In each iteration, the function splits the series to two parts, fits a separate line (using `series_fit_line()`) to each part, and calculates the total r-square. The best split is the one that maximized r-square; the function returns its parameters:

Parameter	Description
<code>rsquare</code>	R-square ↗ is standard measure of the fit quality. It's a number in the range [0-1], where 1 - is the best possible fit, and 0 means the data is unordered and don't fit any line.
<code>split_idx</code>	The index of breaking point to two segments (zero-based).
<code>variance</code>	Variance of the input data.
<code>rvariance</code>	Residual variance, which is the variance between the input data values the approximated ones (by the two line segments).
<code>line_fit</code>	Numerical array holding a series of values of the best fitted line. The series length is equal to the length of the input array. It's mainly used for charting.
<code>right_rsquare</code>	R-square of the line on the right side of the split, see <code>series_fit_line()</code> .
<code>right_slope</code>	Slope of the right approximated line (of the form $y=ax+b$).
<code>right_interception</code>	Interception of the approximated left line (b from $y=ax+b$).
<code>right_variance</code>	Variance of the input data on the right side of the split.
<code>right_rvariance</code>	Residual variance of the input data on the right side of the split.
<code>left_rsquare</code>	R-square of the line on the left side of the split, see <code>series_fit_line()</code> .
<code>left_slope</code>	Slope of the left approximated line (of the form $y=ax+b$).
<code>left_interception</code>	Interception of the approximated left line (of the form $y=ax+b$).
<code>left_variance</code>	Variance of the input data on the left side of the split.
<code>left_rvariance</code>	Residual variance of the input data on the left side of the split.

ⓘ Note

This function returns multiple columns and so cannot be used as an argument for another function.

Syntax

```
project series_fit_2lines(series)
```

- Will return all mentioned above columns with the following names:
series_fit_2lines_x_rsquare, series_fit_2lines_x_split_idx etc.

```
project (rs, si, v)=series_fit_2lines(series)
```

- Will return the following columns: rs (r-square), si (split index), v (variance) and the rest will look like series_fit_2lines_x_rvariance, series_fit_2lines_x_line_fit and etc.

```
extend (rs, si, v)=series_fit_2lines(series)
```

- Will return only: rs (r-square), si (split index) and v (variance).

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values.

💡 Tip

The most convenient way of using this function is applying it to the results of make-series operator.

Examples

Run the query

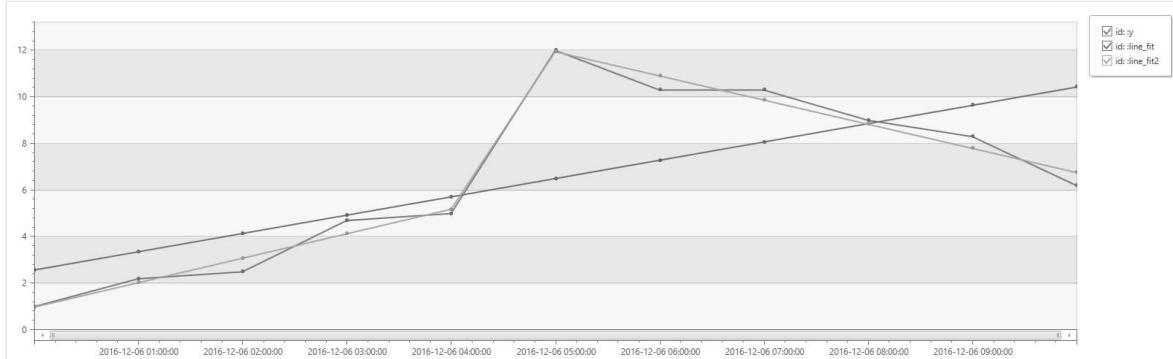
Kusto

```
print
    id=' ',
    x=range(bin(now(), 1h) - 11h, bin(now(), 1h), 1h),
```

```

y=dynamic([1, 2.2, 2.5, 4.7, 5.0, 12, 10.3, 10.3, 9, 8.3, 6.2])
| extend
    (Slope, Interception, RSquare, Variance, RVariance,
LineFit)=series_fit_line(y),
    (RSquare2, SplitIdx, Variance2, RVariance2,
LineFit2)=series_fit_2lines(y)
| project id, x, y, LineFit, LineFit2
| render timechart

```



Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_fit_2lines_dynamic()

Article • 02/01/2023

Applies two segments linear regression on a series, returning a dynamic object.

Takes an expression containing dynamic numerical array as input and applies two segments linear regression ↗ in order to identify and quantify trend changes in a series. The function iterates on the series indexes. In each iteration, it splits the series to two parts, and fits a separate line using `series_fit_line()` or `series_fit_line_dynamic()`. The function fits the lines to each of the two parts, and calculates the total R-squared value. The best split is the one that maximizes R-squared. The function returns its parameters in dynamic value with the following content:

- `rsquare`: R-squared ↗ is a standard measure of the fit quality. It's a number in the range of [0-1], where 1 is the best possible fit, and 0 means the data is unordered and don't fit any line.
- `split_idx`: the index of breaking point to two segments (zero-based).
- `variance`: variance of the input data.
- `rvariance`: residual variance that is the variance between the input data values the approximated ones (by the two line segments).
- `line_fit`: numerical array holding a series of values of the best fitted line. The series length is equal to the length of the input array. It's used for charting.
- `right.rsquare`: r-square of the line on the right side of the split, see `series_fit_line()` or `series_fit_line_dynamic()`.
- `right.slope`: slope of the right approximated line (of the form $y=ax+b$).
- `right.interception`: interception of the approximated left line (b from $y=ax+b$).
- `right.rvariance`: residual variance of the input data on the right side of the split.
- `left.rsquare`: r-square of the line on the left side of the split, see `[series_fit_line()]`. (series-fit-linefunction.md) or `series_fit_line_dynamic()`.
- `left.slope`: slope of the left approximated line (of the form $y=ax+b$).
- `left.interception`: interception of the approximated left line (of the form $y=ax+b$).
- `left.variance`: variance of the input data on the left side of the split.
- `left.rvariance`: residual variance of the input data on the left side of the split.

This operator is similar to `series_fit_2lines`. Unlike `series-fit-2lines`, it returns a dynamic bag.

Syntax

```
series_fit_2lines_dynamic(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values.

💡 Tip

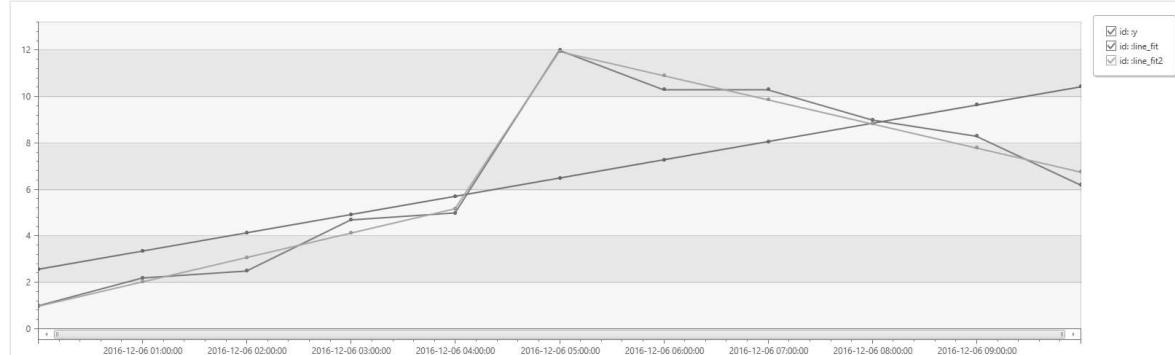
The most convenient way of using this function is applying it to the results of the **make-series** operator.

Example

Run the query

Kusto

```
print
    id=' ',
    x=range(bin(now(), 1h) - 11h, bin(now(), 1h), 1h),
    y=dynamic([1, 2.2, 2.5, 4.7, 5.0, 12, 10.3, 10.3, 9, 8.3, 6.2])
| extend
    LineFit=series_fit_line_dynamic(y).line_fit,
    LineFit2=series_fit_2lines_dynamic(y).line_fit
| project id, x, y, LineFit, LineFit2
| render timechart
```



Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_fit_line()

Article • 02/01/2023

Applies linear regression on a series, returning multiple columns.

Takes an expression containing dynamic numerical array as input and does linear regression[↗] to find the line that best fits it. This function should be used on time series arrays, fitting the output of make-series operator. The function generates the following columns:

- `rsquare`: r-square[↗] is a standard measure of the fit quality. The value's a number in the range [0-1], where 1 - is the best possible fit, and 0 means the data is unordered and doesn't fit any line.
- `slope`: Slope of the approximated line ("a" from $y=ax+b$).
- `variance`: Variance of the input data.
- `rvariance`: Residual variance that is the variance between the input data values the approximated ones.
- `interception`: Interception of the approximated line ("b" from $y=ax+b$).
- `line_fit`: Numerical array holding a series of values of the best fitted line. The series length is equal to the length of the input array. The value's used for charting.

Syntax

```
series_fit_line(series)
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.

💡 Tip

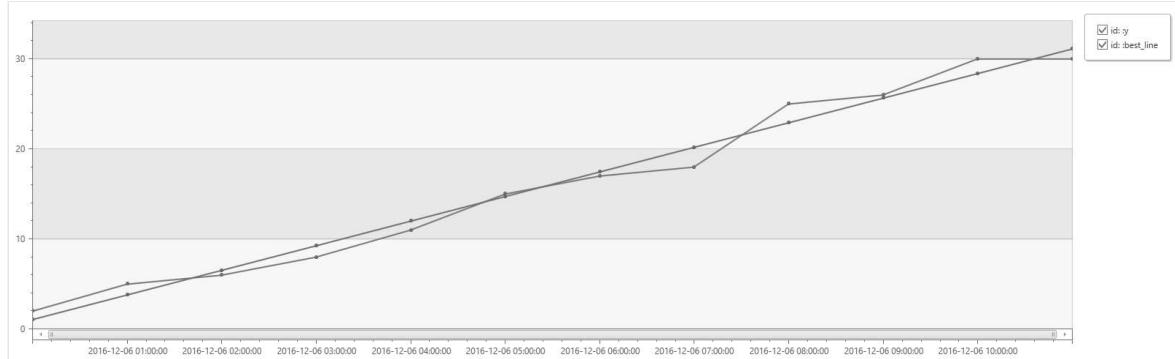
The most convenient way of using this function is to apply it to the results of `make-series` operator.

Examples

Run the query

Kusto

```
print
    id=' ',
    x=range(bin(now(), 1h) - 11h, bin(now(), 1h), 1h),
    y=dynamic([2, 5, 6, 8, 11, 15, 17, 18, 25, 26, 30, 30])
| extend (RSquare, Slope, Variance, RVariance, Interception,
LineFit)=series_fit_line(y)
| render timechart
```



RSquare	Slope	Variance	RVariance	Interception	LineFit
0.982	2.730	98.628	1.686	-1.666	1.064, 3.7945, 6.526, 9.256, 11.987, 14.718, 17.449, 20.180, 22.910, 25.641, 28.371, 31.102

Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_fit_line_dynamic()

Article • 02/01/2023

Applies linear regression on a series, returning dynamic object.

Takes an expression containing dynamic numerical array as input, and does linear regression[↗] to find the line that best fits it. This function should be used on time series arrays, fitting the output of make-series operator. It generates a dynamic value with the following content:

- `rsquare`: r-square[↗] is a standard measure of the fit quality. It's a number in the range [0-1], where 1 is the best possible fit, and 0 means the data is unordered and doesn't fit any line
- `slope`: Slope of the approximated line (the a -value from $y=ax+b$)
- `variance`: Variance of the input data
- `rvariance`: Residual variance that is the variance between the input data values and the approximated ones.
- `interception`: Interception of the approximated line (the b -value from $y=ax+b$)
- `line_fit`: Numerical array containing a series of values of the best fit line. The series length is equal to the length of the input array. It's used mainly for charting.

This operator is similar to `series_fit_line`, but unlike `series-fit-line` it returns a dynamic bag.

Syntax

```
series_fit_line_dynamic( series )
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.

💡 Tip

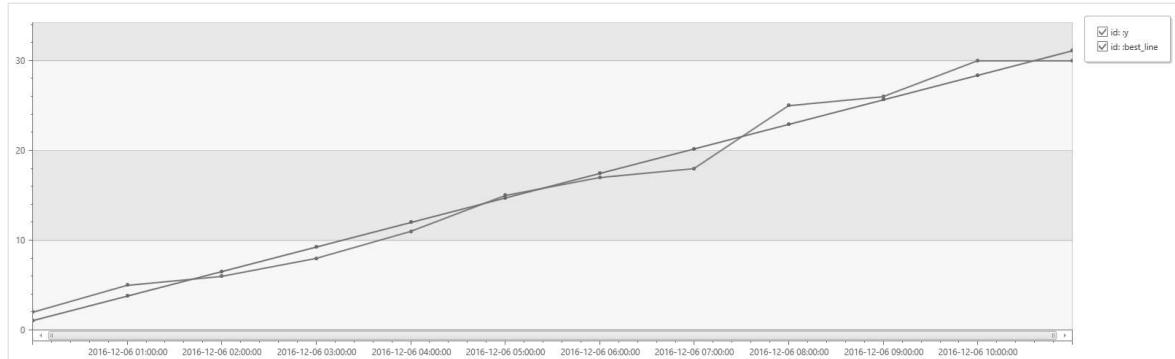
The most convenient way of using this function is by applying it to the results of `make-series` operator.

Examples

Run the query

Kusto

```
print
    id=' ',
    x=range(bin(now(), 1h) - 11h, bin(now(), 1h), 1h),
    y=dynamic([2, 5, 6, 8, 11, 15, 17, 18, 25, 26, 30, 30])
| extend fit=series_fit_line_dynamic(y)
| extend
    RSquare=fit.rsquare,
    Slope=fit.slope,
    Variance=fit.variance,
    RVariance=fit.rvariance,
    Interception=fit.interception,
    LineFit=fit.line_fit
| render timechart
```



RSquare	Slope	Variance	RVariance	Interception	LineFit
0.982	2.730	98.628	1.686	-1.666	1.064, 3.7945, 6.526, 9.256, 11.987, 14.718, 17.449, 20.180, 22.910, 25.641, 28.371, 31.102

Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_fit_poly()

Article • 01/31/2023

Applies a polynomial regression from an independent variable (*x_series*) to a dependent variable (*y_series*). This function takes a table containing multiple series (dynamic numerical arrays) and generates the best fit high-order polynomial for each series using polynomial regression ↴.

💡 Tip

- For linear regression of an evenly spaced series, as created by **make-series operator**, use the simpler function `series_fit_line()`. See [Example 2](#).
- If *x_series* is supplied, and the regression is done for a high degree, consider normalizing to the [0-1] range. See [Example 3](#).
- If *x_series* is of datetime type, it must be converted to double and normalized. See [Example 3](#).
- For reference implementation of polynomial regression using inline Python, see `series_fit_poly_fl()`.

Syntax

```
T | extend series_fit_poly(y_series [, x_series , degree ])
```

Parameters

Name	Type	Required	Description
<i>y_series</i>	dynamic	✓	An array of numeric values containing the dependent variable ↴.
<i>x_series</i>	dynamic		An array of numeric values containing the independent variable ↴. Required only for unevenly spaced series ↴. If not specified, it's set to a default value of [1, 2, ..., length(<i>y_series</i>)].
<i>degree</i>			The required order of the polynomial to fit. For example, 1 for linear regression, 2 for quadratic regression, and so on. Defaults to 1, which indicates linear regression.

Returns

The `series_fit_poly()` function returns the following columns:

- `rsquare`: r-square[↗] is a standard measure of the fit quality. The value's a number in the range [0-1], where 1 - is the best possible fit, and 0 means the data is unordered and doesn't fit any line.
- `coefficients`: Numerical array holding the coefficients of the best fitted polynomial with the given degree, ordered from the highest power coefficient to the lowest.
- `variance`: Variance of the dependent variable (`y_series`).
- `rvariance`: Residual variance that is the variance between the input data values the approximated ones.
- `poly_fit`: Numerical array holding a series of values of the best fitted polynomial. The series length is equal to the length of the dependent variable (`y_series`). The value's used for charting.

Examples

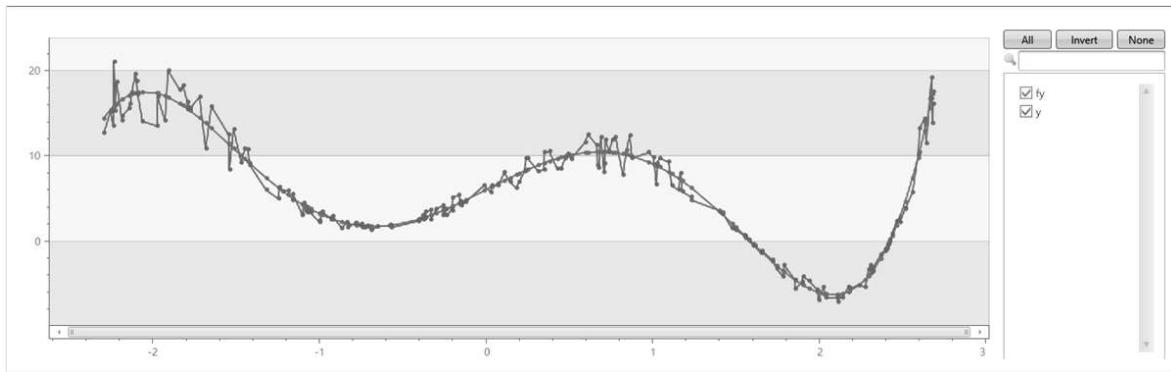
Example 1

A fifth order polynomial with noise on x & y axes:

Run the query

Kusto

```
range x from 1 to 200 step 1
| project x = rand()*5 - 2.3
| extend y = pow(x, 5)-8*pow(x, 3)+10*x+6
| extend y = y + (rand() - 0.5)*0.5*y
| summarize x=make_list(x), y=make_list(y)
| extend series_fit_poly(y, x, 5)
| project-rename fy=series_fit_poly_y_poly_fit,
coeff=series_fit_poly_y_coefficients
| fork (project x, y, fy) (project-away x, y, fy)
| render linechart
```



GenericResult_2				
series_fit_poly_y_rs...	coeff	series_fit_poly_y_va...	series_fit_poly_y_rv...	
0.965150555646126	[1.0295150145250205, -0.09278808879446121, -8.0851056489948121, 0.309264285812111, 9.987337152372735,	44.7735387815632	1.56033294829411	

coeff

Name	JPath
▼ [6]	
1.03	[0]
-0.09	[1]
-8.09	[2]
0.31	[3]
9.99	[4]
5.95	[5]

coeff

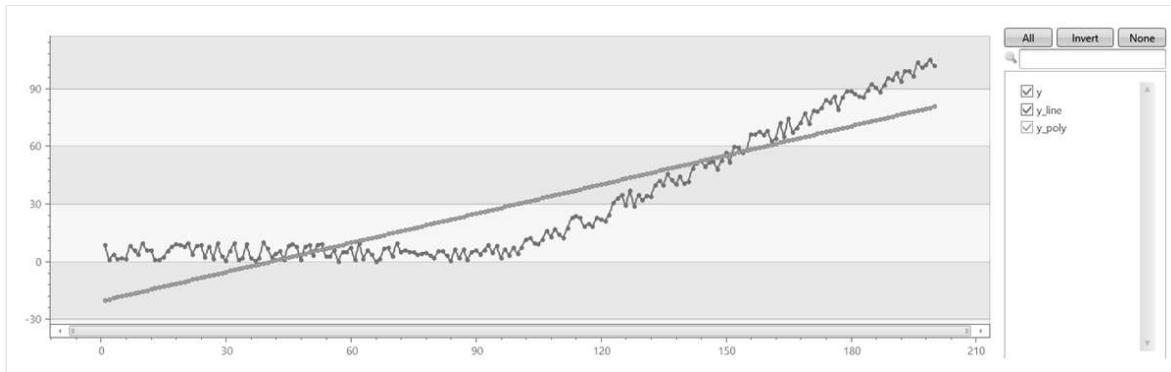
```
[  
1.0295150145250205,  
-0.09278808879446121,  
-8.0851056489948121,  
0.309264285812111,  
9.987337152372735,  
5.9535426896127612  
]
```

Example 2

Verify that `series_fit_poly` with degree=1 matches `series_fit_line`:

Kusto

```
demo_series1
| extend series_fit_line(y)
| extend series_fit_poly(y)
| project-rename y_line = series_fit_line_y_line_fit, y_poly =
series_fit_poly_y_poly_fit
| fork (project x, y, y_line, y_poly) (project-away id, x, y, y_line,
y_poly)
| render linechart with(xcolumn=x, ycolumns=y, y_line, y_poly)
```



GenericResult_2									
series_fit_line_y_rsq...	series_fit_line_y_slo...	series_fit_line_y_var...	series_fit_line_y_rva...	series_fit_line_y_int...	series_fit_poly_y_rs...	series_fit_poly_y_co...	series_fit_poly_y_va...	series_fit_poly_y_v...	series_fit_poly_y_rv...
0.796772368148283	0.50546507398455	1074.21904502052	218.310992609534	-20.5656868292799	0.796772368148283	[0.5054650739845... -20.56568682927...]	1074.21904502052	218.310992609534	

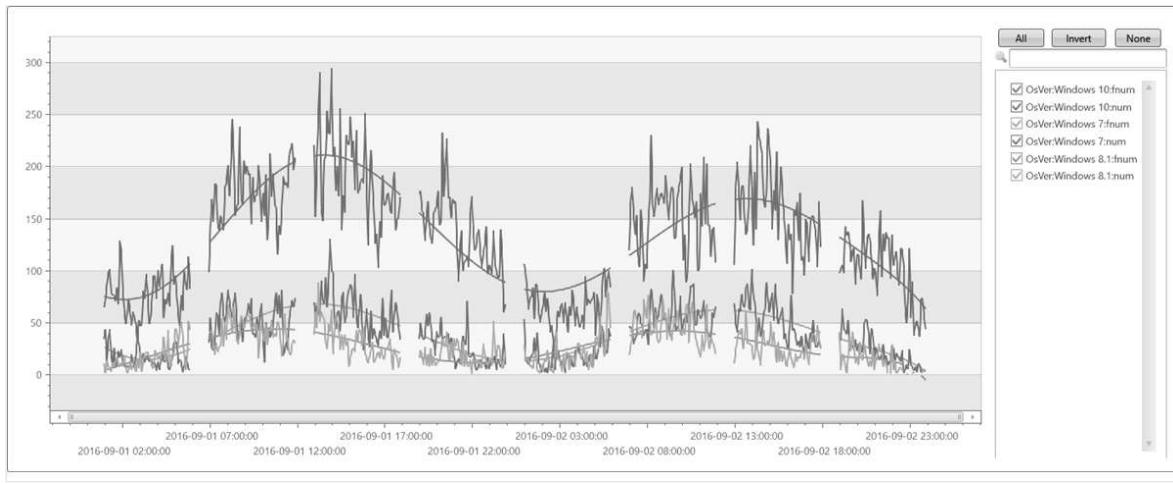
Example 3

Irregular (unevenly spaced) time series:

Run the query

Kusto

```
//  
// x-axis must be normalized to the range [0-1] if either degree is  
// relatively big (>= 5) or original x range is big.  
// so if x is a time axis it must be normalized as conversion of timestamp  
// to long generate huge numbers (number of 100 nano-sec ticks from 1/1/1970)  
//  
// Normalization: x_norm = (x - min(x))/(max(x) - min(x))  
//  
irregular_ts  
| extend series_stats(series_add(TimeStamp, 0))  
// extract min/max of time axis as doubles  
| extend x = series_divide(series_subtract(TimeStamp, series_stats__min),  
series_stats__max-series_stats__min) // normalize time axis to [0-1] range  
| extend series_fit_poly(num, x, 8)  
| project-rename fnum=series_fit_poly_num_poly_fit  
| render timechart with(ycolumns=num, fnum)
```



Feedback

Was this page helpful?

Yes

No

Provide product feedback | Get help at Microsoft Q&A

series_floor()

Article • 01/31/2023

Calculates the element-wise floor function of the numeric series input.

Syntax

```
series_floor(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values on which the floor function is applied.

Returns

Dynamic array of the calculated floor function. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
print s = dynamic([-1.5,1,2.5])  
| extend s_floor = series_floor(s)
```

Output

s	s_floor
[-1.5,1,2.5]	[-2.0,1.0,2.0]

Feedback



Was this page helpful?

Provide product feedback  | Get help at Microsoft Q&A

series_greater()

Article • 01/31/2023

Calculates the element-wise greater (`>`) logic operation of two numeric series inputs.

Syntax

```
series_greater(series1, series2)
```

Parameters

Name	Type	Required	Description
<code>series1</code> , <code>series2</code>	dynamic	✓	The arrays of numeric values to be element-wise compared.

Returns

Dynamic array of booleans containing the calculated element-wise greater logic operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
print s1 = dynamic([1,2,4]), s2 = dynamic([4,2,1])
| extend s1_greater_s2 = series_greater(s1, s2)
```

Output

s1	s2	s1_greater_s2
[1,2,4]	[4,2,1]	[false,false,true]

See also

For entire series statistics comparisons, see:

- `series_stats()`
 - `series_stats_dynamic()`
-

Feedback

Was this page helpful?



Yes



No

Provide product feedback | Get help at Microsoft Q&A

series_greater_equals()

Article • 01/31/2023

Calculates the element-wise greater or equals (`>=`) logic operation of two numeric series inputs.

Syntax

```
series_greater_equals(series1, series2)
```

Parameters

Name	Type	Required	Description
<code>series1</code> , <code>series2</code>	dynamic	✓	The arrays of numeric values to be element-wise compared.

Returns

Dynamic array of booleans containing the calculated element-wise greater or equal logic operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
print s1 = dynamic([1,2,4]), s2 = dynamic([4,2,1])
| extend s1_greater_equals_s2 = series_greater_equals(s1, s2)
```

Output

s1	s2	s1_greater_equals_s2
[1,2,4]	[4,2,1]	[false,true,true]

See also

For entire series statistics comparisons, see:

- `series_stats()`
 - `series_stats_dynamic()`
-

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_ifft()

Article • 01/31/2023

Applies the Inverse Fast Fourier Transform (IFFT) on a series.

The `series_ifft()` function takes a series of complex numbers in the frequency domain and transforms it back to the time/spatial domain using the Fast Fourier Transform[↗]. This function is the complementary function of `series_fft`. Commonly the original series is transformed to the frequency domain for spectral processing and then back to the time/spatial domain.

Syntax

```
series_ifft(fft_real [, fft_imaginary])
```

Parameters

Name	Type	Required	Description
<i>fft_real</i>	dynamic	✓	An array of numeric values representing the real component of the series to transform.
<i>fft_imaginary</i>	dynamic		An array of numeric values representing the imaginary component of the series. This parameter should be specified only if the input series contains complex numbers.

Returns

The function returns the complex inverse fft in two series. The first series for the real component and the second one for the imaginary component.

Example

See `series_fft`

Feedback



Was this page helpful?

Provide product feedback | Get help at Microsoft Q&A

series_iir()

Article • 02/22/2023

Applies an Infinite Impulse Response filter on a series.

The function takes an expression containing dynamic numerical array as input, and applies an Infinite Impulse Response[↗] filter. By specifying the filter coefficients, you can use the function to:

- calculate the cumulative sum of the series
- apply smoothing operations
- apply various high-pass[↗], band-pass[↗], and low-pass[↗] filters

The function takes as input the column containing the dynamic array and two static dynamic arrays of the filter's *denominators* and *numerators* coefficients, and applies the filter on the column. It outputs a new dynamic array column, containing the filtered output.

Syntax

```
series_iir(series , numerators , denominators )
```

Parameters

Name	Type	Required	Description
<i>series</i>	dynamic	✓	An array of numeric values, typically the resulting output of make-series or make_list operators.
<i>numerators</i>	dynamic	✓	An array of numeric values, containing the numerator coefficients of the filter.
<i>denominators</i>	dynamic	✓	An array of numeric values, containing the denominator coefficients of the filter.

ⓘ Important

The first element of *a* (that is, *a[0]*) mustn't be zero, to avoid division by 0. See the following formula.

The filter's recursive formula

- Consider an input array X, and coefficients arrays a and b of lengths n_a and n_b respectively. The transfer function of the filter that will generate the output array Y, is defined by:

$$Y_i = a_0^{-1}(b_0X_i + b_1X_{i-1} + \dots + b_{n_b-1}X_{i-n_b-1} - a_1Y_{i-1} - a_2Y_{i-2} - \dots - a_{n_a-1}Y_{i-n_a-1})$$

Example

Calculate a cumulative sum. Use the iir filter with coefficients *denominators*=[1,-1] and *numerators*=[1]:

Run the query

Kusto

```
let x = range(1.0, 10, 1);
print x=x, y = series_iir(x, dynamic([1]), dynamic([1,-1]))
| mv-expand x, y
```

Output

x	y
1.0	1.0
2.0	3.0
3.0	6.0
4.0	10.0

Here's how to wrap it in a function:

Run the query

Kusto

```
let vector_sum=(x: dynamic) {
    let y=array_length(x) - 1;
    todouble(series_iir(x, dynamic([1]), dynamic([1, -1]))[y])
};

print d=dynamic([0, 1, 2, 3, 4])
| extend dd=vector_sum(d)
```

Output

d	dd
[0,1,2,3,4]	10

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_less()

Article • 01/31/2023

Calculates the element-wise less (`<`) logic operation of two numeric series inputs.

Syntax

```
series_less(series1, series2)
```

Parameters

Name	Type	Required	Description
<code>series1</code> , <code>series2</code>	dynamic	✓	The arrays of numeric values to be element-wise compared.

Returns

Dynamic array of booleans containing the calculated element-wise less logic operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
print s1 = dynamic([1,2,4]), s2 = dynamic([4,2,1])
| extend s1_less_s2 = series_less(s1, s2)
```

Output

s1	s2	s1_less_s2
[1,2,4]	[4,2,1]	[true,false,false]

See also

For entire series statistics comparisons, see:

- `series_stats()`
 - `series_stats_dynamic()`
-

Feedback

Was this page helpful?



Yes



No

Provide product feedback | Get help at Microsoft Q&A

series_less_equals()

Article • 01/31/2023

Calculates the element-wise less or equal (`<=`) logic operation of two numeric series inputs.

Syntax

```
series_less_equals(series1, series2)
```

Parameters

Name	Type	Required	Description
<code>series1</code> , <code>series2</code>	dynamic	✓	The arrays of numeric values to be element-wise compared.

Returns

Dynamic array of booleans containing the calculated element-wise less or equal logic operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
print s1 = dynamic([1,2,4]), s2 = dynamic([4,2,1])
| extend s1_less_equals_s2 = series_less_equals(s1, s2)
```

Output

s1	s2	s1_less_equals_s2
[1,2,4]	[4,2,1]	[true,true,false]

See also

For entire series statistics comparisons, see:

- `series_stats()`
 - `series_stats_dynamic()`
-

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_log()

Article • 01/31/2023

Calculates the element-wise natural logarithm function (base-e) of the numeric series input.

Syntax

```
series_log(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values on which the natural logarithm function is applied.

Returns

Dynamic array of the calculated natural logarithm function. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
print s = dynamic([1,2,3])  
| extend s_log = series_log(s)
```

Output

s	s_log
[1,2,3]	[0.0,0.69314718055994529,1.0986122886681098]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_multiply()

Article • 01/31/2023

Calculates the element-wise multiplication of two numeric series inputs.

Syntax

```
series_multiply(series1, series2)
```

Parameters

Name	Type	Required	Description
series1	dynamic	✓	The arrays of numeric values to be element-wise multiplied.
series2			

Returns

Dynamic array of calculated element-wise multiplication operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
range x from 1 to 3 step 1
| extend y = x * 2
| extend z = y * 2
| project s1 = pack_array(x,y,z), s2 = pack_array(z, y, x)
| extend s1_multiply_s2 = series_multiply(s1, s2)
```

Output

s1	s2	s1_multiply_s2
[1,2,4]	[4,2,1]	[4,4,4]

s1	s2	s1_multiply_s2
[2,4,8]	[8,4,2]	[16,16,16]
[3,6,12]	[12,6,3]	[36,36,36]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_not_equals()

Article • 02/06/2023

Calculates the element-wise not equals (`!=`) logic operation of two numeric series inputs.

Syntax

```
series_not_equals(series1, series2)
```

Parameters

Name	Type	Required	Description
<code>series1</code> , <code>series2</code>	dynamic	✓	The arrays of numeric values to be element-wise compared.

Returns

Dynamic array of booleans containing the calculated element-wise not equal logic operation between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` element value.

Example

Run the query

Kusto

```
print s1 = dynamic([1,2,4]), s2 = dynamic([4,2,1])
| extend s1_not_equals_s2 = series_not_equals(s1, s2)
```

Output

s1	s2	s1_not_equals_s2
[1,2,4]	[4,2,1]	[true,false,true]

See also

For entire series statistics comparisons, see:

- `series_stats()`
 - `series_stats_dynamic()`
-

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_outliers()

Article • 02/06/2023

Scores anomaly points in a series.

The function takes an expression with a dynamic numerical array as input, and generates a dynamic numeric array of the same length. Each value of the array indicates a score of a possible anomaly, using "Tukey's test" ↗. A value greater than 1.5 in the same element of the input indicates a rise or decline anomaly. A value less than -1.5, indicates a decline anomaly.

Syntax

```
series_outliers(series [, kind] [, ignore_val] [, min_percentile] [, max_percentile])
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>kind</code>	string		The algorithm to use for outlier detection. The supported options are <code>"tukey"</code> , which is traditional "Tukey", and <code>"ctukey"</code> , which is custom "Tukey". The default is <code>"ctukey"</code> .
<code>ignore_val</code>	int, long, or real		A numeric value indicating the missing values in the series. The default is <code>double(null)</code> . The score of nulls and ignore values is set to <code>0</code> .
<code>min_percentile</code>	int, long, or real		The minimum percentile to use to calculate the normal inter-quantile range. The default is 10. The value must be in the range <code>[2.0, 98.0]</code> . This parameter is only relevant for the <code>"ctukey"</code> <i>kind</i> .
<code>max_percentile</code>	int, long, or real		The maximum percentile to use to calculate the normal inter-quantile range. The default is 90. The value must be in the range <code>[2.0, 98.0]</code> . This parameter is only relevant for the <code>"ctukey"</code> <i>kind</i> .

The following table describes differences between `"tukey"` and `"ctukey"`:

Algorithm	Default quantile range	Supports custom quantile range
-----------	------------------------	--------------------------------

Algorithm	Default quantile range	Supports custom quantile range
"tukey"	25% / 75%	No
"ctukey"	10% / 90%	Yes

💡 Tip

The best way to use this function is to apply it to the results of the `make-series` operator.

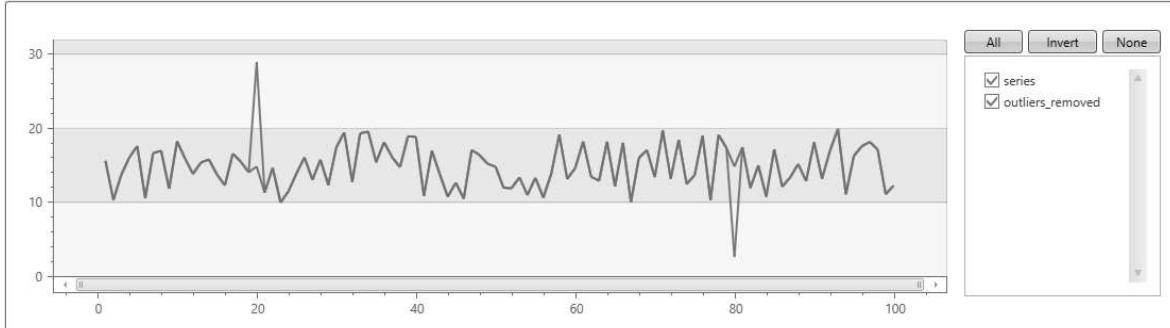
Example

A time series with some noise creates outliers. If you would like to replace those outliers (noise) with the average value, use `series_outliers()` to detect the outliers, and then replace them.

[Run the query](#)

Kusto

```
range x from 1 to 100 step 1
| extend y=iff(x == 20 or x == 80, 10 * rand() + 10 + (50 - x) / 2, 10 * rand() + 10) // generate a sample series with outliers at x=20 and x=80
| summarize x=make_list(x), series=make_list(y)
| extend series_stats(series), outliers=series_outliers(series)
| mv-expand x to typeof(long), series to typeof(double), outliers to typeof(double)
| project
    x,
    series,
    outliers_removed=iff(outliers > 1.5 or outliers < -1.5,
series_stats_series_avg, series) // replace outliers with the average
| render linechart
```



Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_pearson_correlation()

Article • 02/06/2023

Calculates the pearson correlation coefficient of two numeric series inputs.

See: Pearson correlation coefficient [↗](#).

Syntax

```
series_pearson_correlation(series1, series2)
```

Parameters

Name	Type	Required	Description
series1, series2	dynamic	✓	The arrays of numeric values for calculating the correlation coefficient.

Returns

The calculated Pearson correlation coefficient between the two inputs. Any non-numeric element or non-existing element (arrays of different sizes) yields a `null` result.

Example

Run the query

Kusto

```
range s1 from 1 to 5 step 1
| extend s2 = 2 * s1 // Perfect correlation
| summarize s1 = make_list(s1), s2 = make_list(s2)
| extend correlation_coefficient = series_pearson_correlation(s1, s2)
```

Output

s1	s2	correlation_coefficient
[1,2,3,4,5]	[2,4,6,8,10]	1

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_periods_detect()

Article • 02/06/2023

Finds the most significant periods that exist in a time series.

Syntax

```
series_periods_detect(series, min_period, max_period, num_periods)
```

Parameters

Name	Type	Required	Description
<i>series</i>	dynamic	✓	An array of numeric values, typically the resulting output of the make-series or make_list operators.
<i>min_period</i>	real	✓	The minimal period for which to search.
<i>max_period</i>	real	✓	The maximal period for which to search.
<i>num_periods</i>	long	✓	The maximum required number of periods. This number will be the length of the output dynamic arrays.

ⓘ Important

- The algorithm can detect periods containing at least 4 points and at most half of the series length.
- Set the *min_period* a little below and *max_period* a little above the periods you expect to find in the time series. For example, if you have an hourly aggregated signal, and you look for both daily and weekly periods (24 and 168 hours respectively), you can set *min_period*=0.8*24, *max_period*=1.2*168, and leave 20% margins around these periods.
- The input time series must be regular. That is, aggregated in constant bins, which is always the case if it has been created using make-series. Otherwise, the output is meaningless.

Returns

The function outputs a table with two columns:

- *periods*: A dynamic array containing the periods that have been found, in units of the bin size, ordered by their scores.
- *scores*: A dynamic array containing values between 0 and 1. Each array measures the significance of a period in its respective position in the *periods* array.

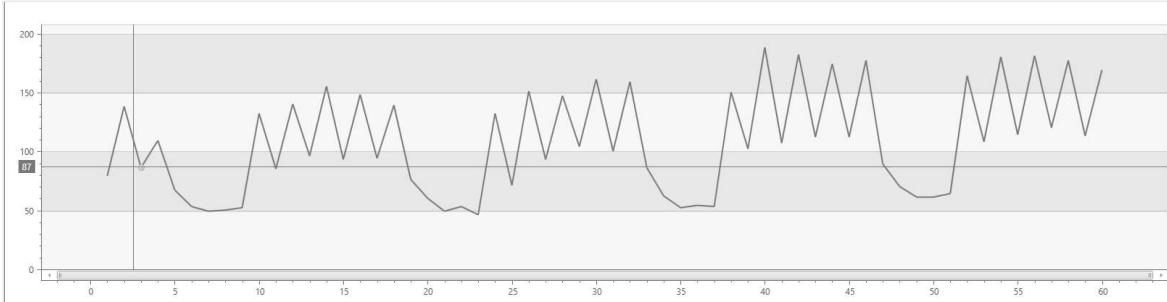
Example

The following query embeds a snapshot of a month of an application's traffic, aggregated twice a day. The bin size is 12 hours.

Run the query

Kusto

```
print y=dynamic([80, 139, 87, 110, 68, 54, 50, 51, 53, 133, 86, 141, 97, 156, 94, 149, 95, 140, 77, 61, 50, 54, 47, 133, 72, 152, 94, 148, 105, 162, 101, 160, 87, 63, 53, 55, 54, 151, 103, 189, 108, 183, 113, 175, 113, 178, 90, 71, 62, 62, 65, 165, 109, 181, 115, 182, 121, 178, 114, 170]) | project x=range(1, array_length(y), 1), y | render linechart
```



Running `series_periods_detect()` on this series, results in the weekly period, 14 points long.

Run the query

Kusto

```
print y=dynamic([80, 139, 87, 110, 68, 54, 50, 51, 53, 133, 86, 141, 97, 156, 94, 149, 95, 140, 77, 61, 50, 54, 47, 133, 72, 152, 94, 148, 105, 162, 101, 160, 87, 63, 53, 55, 54, 151, 103, 189, 108, 183, 113, 175, 113, 178, 90, 71, 62, 62, 65, 165, 109, 181, 115, 182, 121, 178, 114, 170]) | project x=range(1, array_length(y), 1), y | project series_periods_detect(y, 0.0, 50.0, 2)
```

Output

series_periods_detect_y_periods	series_periods_detect_y_periods_scores
[14.0, 0.0]	[0.84, 0.0]

ⓘ Note

The daily period that can be also seen in the chart wasn't found because the sampling is too coarse (12h bin size), so a daily period of 2 bins is below the minimum period size of 4 points, required by the algorithm.

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_periods_validate()

Article • 05/09/2023

Checks whether a time series contains periodic patterns of given lengths.

Often a metric measuring the traffic of an application is characterized by a weekly or daily period. This period can be confirmed by running `series_periods_validate()` that checks for a weekly and daily period.

Syntax

```
series_periods_validate(series, period1 [, period2, ... ] )
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values, typically the resulting output of make-series or make_list operators.
<code>period1,</code> <code>period2,</code> etc.	real	✓	The periods to validate in units of the bin size. For example, if the series is in 1h bins, a weekly period is 168 bins. At least one period is required.

ⓘ Important

- The minimal value for each of the `period` parameters is **4** and the maximal is half of the length of the input series. For a `period` argument outside these bounds, the output score will be **0**.
- The input time series must be regular, that is, aggregated in constant bins, and is always the case if it has been created using `make-series`. Otherwise, the output is meaningless.
- The function accepts up to 16 periods to validate.

Returns

The function outputs a table with two columns:

- *periods*: A dynamic array that contains the periods to validate as supplied in the input.
- *scores*: A dynamic array that contains a score between 0 and 1. The score shows the significance of a period in its respective position in the *periods* array.

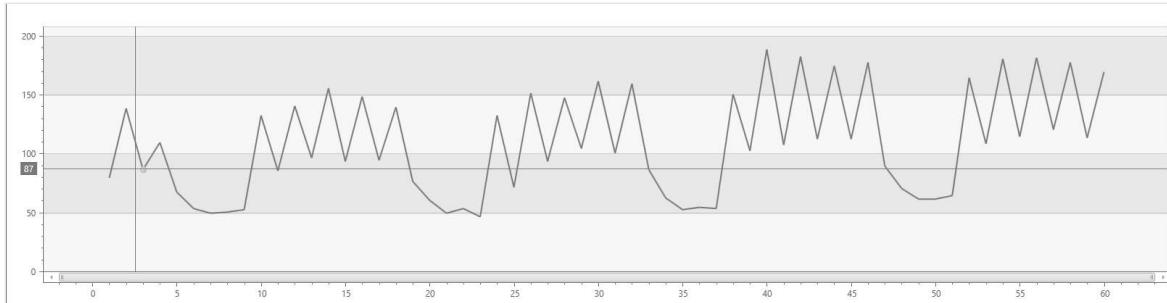
Example

The following query embeds a snapshot of a month of an application's traffic, aggregated twice a day (the bin size is 12 hours).

[Run the query](#)

Kusto

```
print y=dynamic([80, 139, 87, 110, 68, 54, 50, 51, 53, 133, 86, 141, 97,
156, 94, 149, 95, 140, 77, 61, 50, 54, 47, 133, 72, 152, 94, 148, 105, 162,
101, 160, 87, 63, 53, 55, 54, 151, 103, 189, 108, 183, 113, 175, 113, 178,
90, 71, 62, 62, 65, 165, 109, 181, 115, 182, 121, 178, 114, 170])
| project x=range(1, array_length(y), 1), y
| render linechart
```



If you run `series_periods_validate()` on this series to validate a weekly period (14 points long) it results in a high score, and with a 0 score when you validate a five-day period (10 points long).

[Run the query](#)

Kusto

```
print y=dynamic([80, 139, 87, 110, 68, 54, 50, 51, 53, 133, 86, 141, 97,
156, 94, 149, 95, 140, 77, 61, 50, 54, 47, 133, 72, 152, 94, 148, 105, 162,
101, 160, 87, 63, 53, 55, 54, 151, 103, 189, 108, 183, 113, 175, 113, 178,
90, 71, 62, 62, 65, 165, 109, 181, 115, 182, 121, 178, 114, 170])
| project x=range(1, array_length(y), 1), y
| project series_periods_validate(y, 14.0, 10.0)
```

Output

series_periods_validate_y_periods	series_periods_validate_y_scores
[14.0, 10.0]	[0.84, 0.0]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_pow()

Article • 02/06/2023

Calculates the element-wise power of two numeric series inputs.

Syntax

```
series_pow(series1, series2)
```

Parameters

Name	Type	Required	Description
series1	dynamic	✓	Arrays of numeric values. The first array, or base, is element-wise
series2			raised to the power of the second array, or power, into a dynamic array result.

Returns

A dynamic array of calculated element-wise power operation between the two inputs. Any non-numeric element or non-existing element, such as in the case of arrays of different sizes, yields a `null` element value.

Example

Run the query

Kusto

```
print x = dynamic([1, 2, 3, 4]), y=dynamic([1, 2, 3, 0.5])
| extend x_pow_y = series_pow(x, y)
```

Output

x	y	x_pow_y
[1,2,3,4]	[1,2,3,0.5]	[1.0,4.0,27.0,2.0]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_seasonal()

Article • 02/06/2023

Calculates the seasonal component of a series, according to the detected or given seasonal period.

Syntax

```
series_seasonal(series [, period])
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values.
period	int		The number of bins for each seasonal period. This value can be any positive integer. By default, the value is set to -1, which automatically detects the period using the series_periods_detect() with a threshold of 0.7. If seasonality is not detected, the function returns zeros. If a different value is set, it ignores seasonality and returns a series of zeros.

Returns

A dynamic array of the same length as the *series* input that contains the calculated seasonal component of the series. The seasonal component is calculated as the *median* of all the values that correspond to the location of the bin, across the periods.

Examples

Auto detect the period

In the following example, the series' period is automatically detected. The first series' period is detected to be six bins and the second five bins. The third series' period is too short to be detected and returns a series of zeroes. See the next example on how to force the period.

Run the query

Kusto

```
print s=dynamic([2, 5, 3, 4, 3, 2, 1, 2, 3, 4, 3, 2, 1, 2, 3, 4, 3, 2, 1])
| union (print s=dynamic([8, 12, 14, 12, 10, 10, 12, 14, 12, 10, 10, 12, 14, 12, 10, 10, 12, 14, 12, 10]))
| union (print s=dynamic([1, 3, 5, 2, 4, 6, 1, 3, 5, 2, 4, 6]))
| extend s_seasonal = series_seasonal(s)
```

Output

s	s_seasonal
[2,5,3,4,3,2,1,2,3,4,3,2,1,2,3,4,3,2,1,2,3,4,3,2,1]	[1,0,2,0,3,0,4,0,3,0,2,0,1,0,2,0,3,0,4,0,3,0,2,0,1,0,2,0,3,0,4,0,3,0,2,0,1,0]
[8,12,14,12,10,10,12,14,12,10,10,12,14,12,10,10,12,14,12,10]	[10,0,12,0,14,0,12,0,10,0,10,0,12,0,14,0,12,0,10,0,10,0,12,0,14,0,12,0,10,0]
[1,3,5,2,4,6,1,3,5,2,4,6]	[0,0]

Force a period

In this example, the series' period is too short to be detected by series_periods_detect(), so we explicitly force the period to get the seasonal pattern.

Run the query

Kusto

```
print s=dynamic([1, 3, 5, 1, 3, 5, 2, 4, 6])
| union (print s=dynamic([1, 3, 5, 2, 4, 6, 1, 3, 5, 2, 4, 6]))
| extend s_seasonal = series_seasonal(s, 3)
```

Output

s	s_seasonal
[1,3,5,1,3,5,2,4,6]	[1.0,3.0,5.0,1.0,3.0,5.0,1.0,3.0,5.0]
[1,3,5,2,4,6,1,3,5,2,4,6]	[1.5,3.5,5.5,1.5,3.5,5.5,1.5,3.5,5.5,1.5,3.5,5.5]

See also

- [series_periods_detect\(\)](#)
- [series_periods_validate\(\)](#)

Feedback

Was this page helpful?

Provide product feedback [↗](#) | Get help at Microsoft Q&A

series_sign()

Article • 02/06/2023

Calculates the element-wise sign of the numeric series input.

Syntax

```
series_sign(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the sign function is applied.

Returns

A dynamic array of calculated sign function values. -1 for negative, 0 for 0, and 1 for positive. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
  
print arr = dynamic([-6, 0, 8])  
| extend arr_sign = series_sign(arr)
```

Output

arr	arr_sign
[-6,0,8]	[-1,0,1]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_sin()

Article • 02/06/2023

Calculates the element-wise sine of the numeric series input.

Syntax

```
series_sin(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values over which the sine function is applied.

Returns

A dynamic array of calculated sine function values. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
  
print arr = dynamic([-1, 0, 1])  
| extend arr_sin = series_sin(arr)
```

Output

arr	arr_sin
[-6.5,0,8.2]	[-0.8414709848078965,0.0,0.8414709848078965]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_stats()

Article • 03/29/2023

Returns statistics for a numerical series in a table with a column for each statistic.

⚠ Note

This function returns multiple values. If you only need a single value, such as the average, consider using `series_stats_dynamic`.

Syntax

```
... | extend ( Name, ... ) = series_stats ( series [, ignore_nonfinite] )
```

Parameters

Name	Type	Required	Description
<code>Name</code>	string		The column labels for the output table. If not provided, the system will generate them. If you provide a limited number of names, the table will show only those columns.
<code>series</code>	dynamic	✓	An array of numeric values.
<code>ignore_nonfinite</code>	bool		Determines if the calculation includes non-finite values like <code>null</code> , <code>NaN</code> , <code>inf</code> , and so on. The default is <code>false</code> , which will result in <code>null</code> if non-finite values are present.

Returns

A table with a column for each of the statistics displayed in the following table.

Statistic	Description
<code>min</code>	The minimum value in the input array.
<code>min_idx</code>	The first position of the minimum value in the input array.
<code>max</code>	The maximum value in the input array.
<code>max_idx</code>	The first position of the maximum value in the input array.
<code>avg</code>	The average value of the input array.
<code>variance</code>	The sample variance of input array.
<code>stdev</code>	The sample standard deviation of the input array.

Example

Run the query

```
Kusto
print x=dynamic([23, 46, 23, 87, 4, 8, 3, 75, 2, 56, 13, 75, 32, 16, 29])
| project series_stats(x)
```

Output

series_stats_x_min	series_stats_x_min_idx	series_stats_x_max	series_stats_x_max_idx	series_stats_x_avg	series_stats_x_stdev	series_stats_x_variance
2	8	87	3	32.8	28.5036338535483	812.457142857143

Feedback

Was this page helpful?  Yes  No

Provide product feedback  | Get help at Microsoft Q&A

series_stats_dynamic()

Article • 02/06/2023

Returns statistics for a series in a dynamic object.

Syntax

```
series_stats_dynamic(series [, ignore_nonfinite ])
```

Parameters

Name	Type	Required	Description
<code>series</code>	dynamic	✓	An array of numeric values.
<code>ignore_nonfinite</code>	bool		Indicates whether to calculate the statistics while ignoring non-finite values, such as <code>null</code> , <code>Nan</code> , <code>inf</code> , and so on. The default is <code>false</code> , which returns <code>null</code> if non-finite values are present in the array.

Returns

A dynamic property bag object with the following content:

- `min`: The minimum value in the input array.
- `min_idx`: The first position of the minimum value in the input array.
- `max`: The maximum value in the input array.
- `max_idx`: The first position of the maximum value in the input array.
- `avg`: The average value of the input array.
- `variance`: The sample variance of input array.
- `stdev`: The sample standard deviation of the input array.
- `sum`: The sum of the values in the input array.
- `len`: The length of the input array.

Example

Run the query

```
print x=dynamic([23, 46, 23, 87, 4, 8, 3, 75, 2, 56, 13, 75, 32, 16, 29])
| project stats=series_stats_dynamic(x)
```

Output

stats

```
{"min": 2.0, "min_idx": 8, "max": 87.0, "max_idx": 3, "avg": 32.8, "stdev": 28.503633853548269,
"variance": 812.45714285714291, "sum": 492.0, "len": 15}
```

Feedback

Was this page helpful?



Provide product feedback | Get help at Microsoft Q&A

series_subtract()

Article • 02/06/2023

Calculates the element-wise subtraction of two numeric series inputs.

Syntax

```
series_subtract(series1, series2)
```

Parameters

Name	Type	Required	Description
series1,	dynamic	✓	Arrays of numeric values, the second array to be element-wise subtracted from the first array.
series2			

Returns

A dynamic array of calculated element-wise subtract operation between the two inputs. Any non-numeric element or non-existing element, such as in the case of arrays of different sizes, yields a `null` element value.

Example

Run the query

Kusto

```
range x from 1 to 3 step 1
| extend y = x * 2
| extend z = y * 2
| project s1 = pack_array(x,y,z), s2 = pack_array(z, y, x)
| extend s1_subtract_s2 = series_subtract(s1, s2)
```

Output

s1	s2	s1_subtract_s2
[1,2,4]	[4,2,1]	[-3,0,3]

s1	s2	s1_subtract_s2
[2,4,8]	[8,4,2]	[-6,0,6]
[3,6,12]	[12,6,3]	[-9,0,9]

Feedback

Was this page helpful?

 Yes

 No

Provide product feedback  | Get help at Microsoft Q&A

series_tan()

Article • 02/06/2023

Calculates the element-wise tangent of the numeric series input.

Syntax

```
series_tan(series)
```

Parameters

Name	Type	Required	Description
series	dynamic	✓	An array of numeric values on which the tangent function is applied.

Returns

A dynamic array of calculated tangent function values. Any non-numeric element yields a `null` element value.

Example

Run the query

```
Kusto  
  
print arr = dynamic([-1, 0, 1])  
| extend arr_tan = series_tan(arr)
```

Output

arr	arr_tan
[-6.5,0,8.2]	[-1.5574077246549023,0.0,1.5574077246549023]

Feedback